OracleScene 44

SERVING THE ORACLE COMMUNITY

SUMMER 2011

Plus... an extended interview with Oracle Expert Joel Goodman

DEBRA RISES TO THE CHALLENGE

Introducing the new UKOUG Chairman

ALSO IN THIS EDITION

Brand Building
Grant Ronald on the art of 'skinning'

Data Mining
Dig for data treasure with Brendan Tierney

UKOUG
UK ORACLE USER GROUP

Issue 44 £10.00 www.ukoug.org

AN INDEPENDENT PUBLICATION NOT AFFILIATED WITH ORACLE CORPORATION


UKOUG
PARTNER OF THE YEAR
AWARD WINNER 2011/2012
Together...
We made it
Possible!
18+ TOP
INDUSTRY
AWARDS
E-Business Suite Partner
Hyperion Partner
Siebel Partner
PeopleSoft Partner
JD Edwards Partner
ISV Partner
Database Partner
Middleware Partner
Business Intelligence Partner
Innovation Partner
Training Partner
Managed Services Partner
Public Sector Partner
Financial Services Partner
Retail, Manufacturing and Distribution Partner
Communications, Media & Utilities Partner
Services, Transport & Logistics Partner
Small Business/General Business Partner
Nominations now open:
www.ukoug.org/our-partners/ukoug-pya
UKOUG
UK ORACLE USER GROUP

UKOUG
UK ORACLE USER GROUP

OracleScene

ISSUE 44 | Summer 2011

Welcome to Oracle Scene 44. This is our first 'new look' edition of the magazine. We hope you like it. We wanted to design a magazine that better reflected the vibrancy and diversity of the UK Oracle User Community, as well as provide pages of valuable technical advice and guidance.
In this edition, we welcome our new Chairman and Deputy Chairman, Debra Lilley and Lisa Dobson respectively. We also hear about the forthcoming Volunteer's Day and the Partner of the Year Awards (get those applications in folks!). In our Technology section, we've prepared a feast of features including 'Skinning' Skills from Grant Ronald, Mining for Data with Brendan Tierney and an extended interview with Oracle expert, Joel Goodman. And, it doesn't stop there... Cussons talks about OrgPlus in E-Business Suite plus Hyperion is on the ball with comprehensive coverage of their new relationship with Oracle as well as excellent guidance on Hyperion Planning. You can also read about PeopleSoft's plans for the future and why Mogens Nørgaard is hanging some shirts out to dry! Happy reading.

Published by
UK Oracle User Group, User Group House,
591-593 Kingston Road, Wimbledon
London, SW20 8SA
Tel: +44 (0) 20 8545 9670
Fax: +44 (0) 870 9000 335
Email: info@ukoug.org
Web: www.ukoug.org

Editor: Neil Jarvis, editor@ukoug.org.uk
Deputy Editor: Geoff Swaffer
Design: doggARTadams.com
Front Cover: UKOUG Chairman, Debra Lilley, Victoria Centre, Belfast City Centre. Photography by doggARTadams.

UKOUG Directors: Debra Lilley , Chairman (debra.lilley@ukoug.org.uk); Lisa Dobson, Deputy Chairman, (lisa.dobson@ukoug.org.uk); Anne Power, Jeremy Duggan, Carl Dudley, David Kurtz, David Rowntree, Graham Spicer, Robert Stanton, Ronan Miles, Peter Robson, Sue Yates. Directors may be contacted at director@ukoug.org.

IN THIS ISSUE...

RE-DISCOVERING YOUR SENSE OF ADVENTURE

A special message from the new UKOUG Chairman, Debra Lilley

OS: Debra, congratulations on your new position as Chairman of the UKOUG. You must be very proud?
Debra: Yes. I am really very, very proud of UKOUG and I'm very proud to be leading it for the next year. I've been on the board of the UKOUG for eleven years and for nine of those years I've been Deputy Chairman to Ronan Miles. During that time, I hope I've done a lot for the User Group and, what's been important for me, is that I've shown that partners have a lot to offer without engaging in direct marketing. So, I feel that I've earned the respect of people as an individual. When Ronan decided to step down from the role of Chairman, I hoped that the other board members would give me the opportunity to be Chairman. And, I'm very pleased that they did.

OS: Ronan has made an incredible contribution to the UKOUG hasn't he?
Debra: Phenomenal. I don't know how anyone could replicate the contribution he has made. He took over the User Group when it was probably at its strongest. Ronan worked to shore up UKOUG for the future and then, when Oracle began its acquisition policy a few years ago, he did an awful lot to expand and support the communities.

OS: What will your priorities be over the next 12 months?
Debra: My priorities are to bring change to the User Group. The community has changed. What we need to be able to do is alter our model to suit a more diverse community. A lot of our members now aren't Oracle 'wall-to-wall'. For example, some only use Oracle for parts of their business and some are Open Source. So, it is a very different community to what it was a few years ago. Also, the world today is very different and we need to move forward and change our organisation to better meet those needs.

OS: When you bring about change isn't there a danger that you could dilute what the UKOUG represents?
Debra: Nothing will be diluted. No. It will be strengthened. We bring education to users. We create opportunities to help them learn from each other and opportunities to help them get a better return on investment in Oracle software. Different communities need to get that information in different ways. The younger users out there don't necessarily want to meet and to talk to people. They want web content. Some of our communities prefer to meet in a bar after work and talk for a few hours and don't have a particularly rigid agenda. And yet, we still have a large community that wants to have several days of sessions. We just need to be able to adapt to each of those community needs. We also need to adapt to the funding needs. Our membership model is based around corporations and as our community grows the membership doesn't necessarily grow in a number of corporations. For example, people who are part of a new acquisition, their organisation probably has Oracle somewhere so they don't need a new membership. Marketing budgets are often cut or under pressure. Organisations are also merging so exhibition space and sponsorship has been declining. Today, we face many challenges rasing the income we need and we just have to become agile enough to deal with this new environment. It doesn't mean we 'dumb down' what we do. In fact, we do the opposite. We become leaner, more efficient, more responsive and more adventurous in meeting the needs of all of our users.

> "Different communities need information in different ways"

OS: So, I take it you have a pretty busy schedule ahead of you?
Debra: Yes. I'm still a very active speaker in the Oracle community which I love. The User Group is all about educating people and I love being one of those educators. The ACE Director's Programme allows me to speak around the world. I'm just back the EMEA Harmony Conference, in Helsinki. I'll also be attending the ODTUG Kscope Conference, in California. Next month, I hope to tour Latin America. The beginning of July is, of course, a really busy time for UKOUG. We have our Annual General Meeting as well as our Volunteer's Day where we get input from all of our Community Leaders and then our agenda planning day for our main conference in December.

OS: And... of course there we can look forward to this year's Partner of the Year Awards...
Debra: You know it is really important that users can see who are the partners that customers are recommending. We talked to our partner members and they want wanted awards that would be recognised by users. I encourage partners to have a look at the different criteria for the awards, and think which ones are related to them... and get their nominations in. It is a very important part of our year. This will be our fourth year. It's very popular and it is growing. So nominations please!

OS: What's your message for members of the UKOUG?
Debra: I think I'd say to the users, look at the agendas, look at our calendar very closely. There is something there for everybody. I believe our mainstay users know that. What I'd like to say to their managers is that we offer the most cost effective way of delivering training. So, have a look at what's on offer and send your staff.
UKOUG is not just about me, I have the entire Board, staff and volunteers to help me. I'd also like to thank Fujitsu, my employer, for the time they give me to work with UKOUG.
debra.lilley@ukoug.org.uk

Photo Left: UKOUG Chairman, Debra Lilley, on a recent diving vacation in Grand Canaria.

...New Deputy will help drive UKOUG forward.

INTRODUCING LISA DOBSON

I love volunteering with UKOUG. Since joining, in 2005, I have met so many fantastic people and made amazing friendships.

I joined the Board as a Director in 2008 and it has been two years of hard work. It is fun and rewarding, but there is no denying that it does require time and effort in addition to holding down a day job.

Whilst I am hugely excited about taking on this new role, I didn't quite get off to the start I would have wanted. I was unable to attend the March Board meeting in person, due to a rather unfortunate skiing accident that resulted in a fractured leg. So, I had to participate via a Skype video call. It was a rather surreal experience to give my pitch to the other Directors and try to explain why I felt I was the right person for the role via my laptop. I would have much preferred to have been there but it just wasn't possible.

It's been rather a frustrating time since the accident as I haven't been able to get out and about as much as I would have liked, but it has given me plenty of time to get my head around the mammoth tasks ahead. Since being elected, I've been working closely with Debra who has been helping me with the change in addition to her own transition into the role of Chairman. Debra and I are very good friends outside of UKOUG so I know we will work well together.

We've also been doing a lot of work around changes to the organisation as a whole, starting with the boardroom and continuing through to the relationship with the board, the office and our team of volunteers. We need to ensure that the organisation is run effectively and that we continue to meet the needs of our members.

> "...a huge success - a fantastic agenda with great speakers and a packed room"

I'm now well on the way to a full recovery, although I had been very disappointed when I was unable to make it to the 'Back-to-Basics' event, in March, I was still able to make April's 'Exadata' Special Event, in London. The travel arrangements, however, had to be made with military precision. Thanks to some good friends, piggyback rides and a Miracle bag it was a trouble free couple of days! It's events like this that make me very proud to be part of UKOUG.

The members asked for content on Exadata. We responded and the event itself was a huge success - a fantastic agenda with great speakers and a packed room. We were very proud to have been able to pull it all together at quite short notice.

I am very proud to be working with Debra to lead UKOUG through these changes and I ask that you support us by continuing to work with us and the office by providing valuable feedback.

You can also support us by volunteering with UKOUG. We are always looking for people to help out with events, presentations, writing for Oracle Scene and finding other ways in which to share their knowledge and expertise with the community.

lisa.dobson@ukoug.org.uk

Grant Ronald
Oracle Senior Group Product Manager
investigates 'SKINNING'

BRAND BUILDING

BEAUTY THAT'S SKIN DEEP

Did you know that reds and oranges in a restaurant encourage diners to eat quickly and leave? Red also makes food more appealing and influences people to eat more. Now think of the colour schemes associated with McDonald's.

Colour can influence our emotions in a variety of ways but, more importantly, they can be intrinsically linked to a corporate brand.

So, when it comes to building applications that represent your company, the continuation of that branding into your applications is a vital factor. Think about it, probably the first sensory touch point that a customer using your application will have, is the colour.

In this article I explain how application look and feel, and corporate branding can be built into your applications using the skinning feature of Oracle ADF.

The Fusion Skin

Oracle has an army of usability experts responsible for understanding how users interact with applications. This means that the UI controls (ADF Faces) used to build Fusion applications are built with usability and visual aesthetics as a key feature of their design

While you might be happy to fully embrace the look, feel and style of these components as is, you may want the flexibility to apply your own corporate branding, while still utilising the rich set of features associated with ADF Faces UI components.

This ability to define your own look and feel is called, appropriately enough, skinning. So any ADF Faces component can have your own corporate branded "skin" applied to it.

As touched on in earlier articles in this series, ADF Faces UI components render themselves in the browser using HTML and JavaScript as the mark-up language. Given this fact, you may have already guessed that CSS (Cascading Style Sheets) are the mechanism by which standards for UI components are defined. A CSS file contains style directives that are grouped into style classes. Each style class can then be applied, using the appropriate mark-up tag, to the page content to effect the appropriate style change. So, does this mean you can set up a CSS for the look and feel of, say a text field, and then apply it to all your input fields using the StyleClass property of the component? Well you could, but you probably won't get the result you expect! When a component such as af:inputText is rendered at runtime, what is generated is maybe a bit more HTML than you might expect. For af:inputText the component will generate HTML tags for a table row <tr>, which is used to line up the label <label> and the input field <input>. Figure 1 shows the browser mark-up that is generated. If you defined a style class using the StyleClass property of the

TECHNOLOGY

```
<tr id="soc2" class="xlu">
<tr id="it1" class="myStyleClass xlu">
<td class="xu x4z">
<td class="xpu" nowrap="" valign="top">
<input id="it1::content" class="x25" type="text" value="Administration" maxlength="30"
auto:" name="it1"/>
</td>
</tr>
```

Figure 1 - The mark-up generated for a single af:inputText component

component then the style class is applied to the "top level" of the generated HTML. You can see this in the mark-up that myStyleClass has been applied to the "root" of the mark-up for that component.

Furthermore, the selector also allows different styles to be associated with other features such as hover (when the mouse hovers over the field it will automatically change to be Aqua), the font and

```
af|inputText::content{background-color: Silver; color: Navy;}
af|inputText:hover{background-color: Aqua;}
af|inputText::access-key{font-weight: bold; color: Navy;}
af|inputText::label{color: Blue;}
```

Figure 2 - Skinning selectors

That means that if your style class defined a background colour of grey then the whole row would show a grey background because the style class would be applied to the table row <tr>, not the input field.

How Skinning Works

But fear not, the clever people who build Oracle ADF have thought this one through. In Oracle ADF each ADF Faces component has a selector for defining a style.
Each selector has a number of elements (or classes), each of which is associated with a feature of the component. For example, as shown in Figure 2 af:inputText has a selector (appropriately called af|inputText) that has an element called "content" that allows you to specify the background colour and the text colour of the main content of the component; in this case, Silver and Navy respectively. This means that without any changes to your UI components or having to specifically set component properties, you can quickly and easily define a style that will affect all instances of an af:inputText component.

colour of the access key, and the colour of the label associated with the input text field, to name only a few.
So, where are these selectors defined and how will an application know to use them? These selectors would be written in a CSS file and then there is a relatively simple set of configuration files involved in defining which CSS skin file (since you could define many different skins) should be used.
These configuration files are shown in Figure 3 Trinidad-config.xml simply states which skin the application will use.
There are a number of default skins you get with Oracle ADF such as "fusion", "simple" and "blafplus-rich". However, if you want to create and use a new skin called,

Figure 3 - Trinidad-confiig.xml

```
<?xml version="1.0" encoding="windows-1252"?>
<trinidad-config xmlns="http://myfaces.apache.org/trinidad/config">
  <skin-family>mycompany_skin</skin-family>
</trinidad-config>

<?xml version="1.0" encoding="windows-1252" ?>
<skins xmlns="http://myfaces.apache.org/trinidad/skin">
    <!-- mycompany skin extends blafplus-rich.desktop so that changes in the
    fod skin are applied on top of this skin -->
    <skin>
        <id>mycompany_skin.desktop</id>
        <family>mycompany_skin</family>
        <render-kit-id>org.apache.myfaces.trinidad.desktop</render-kit-id>
        <extends>blafplus-rich.desktop</extends>
        <style-sheet-name>skins/mycompany_skin.css</style-sheet-name>
    </skin>
    <!-- FOD skin extends fusion.desktop so that changes in the fod skin are
    applied on top of this skin -->
    <skin>
        <id>fod_fusion_skin.desktop</id>
        <family>fod_fusion_skin</family>
        <render-kit-id>org.apache.myfaces.trinidad.desktop</render-kit-id>
        <extends>fusion.desktop</extends>
        <style-sheet-name>skins/fod_skin.css</style-sheet-name>
    </skin>
</skins>
```

Figure 4 - Defining skins and their CSS files

for example, "mycompany_skin", you would include that skin name into the trinidad-config.xml file. This name references an entry in

Figure 5 - Skinning the icons on a tree control

```
&agent ie {
  af|inputText:disabled::label {color: Red;}
}

&agent mozilla {
  af|inputText:disabled::label {color: Orange;}
}
```

Figure 6 - Defining different styles based on browser

another file, trinidad-skins.xml, in which all the skins your application might use are registered. In this file you can define that your new skin extends any existing skin. This is a useful feature because then you don't have to define a look and feel for every possible component; you can just define styles for those selectors that are different from the base skin. In Figure 4, mycompany-skin extends the "blafplus-rich" skin and points to the CSS file "mycompany_skin.css" which is the file in which the style selectors are defined.

Beyond Colour & Fonts

While colours and fonts are some of the primary use cases for skinning, Oracle ADF has more power to control the branding of an application. Skinning can be used to change the icons and graphics that are used in an Oracle ADF application. For example, the icons that are used to represent the expand/collapse nodes of a tree control can be changed to be something more representative of the data you are displaying in the component, as shown in Figure 5. You can even skin the icons used in layout containers such a panels, accordions and tabs. For example, the af:panelBox component uses icons to define the rounded corners and background of the header bar.
Using skinning you could define different icons to give the frame of the component a different shape.

Dynamic Skinning

The skinning developer also has the ability to dynamically define when skins and skin selectors are used. For example, you might choose to apply a different style to a component if the client is running on a Linux client or maybe use different styles depending on the browser being used. To do this, the "agent" tag is used to define which style should be used. In this example (Figure 6), the label for a disabled field will be different depending on whether the application is being rendered in IE or Mozilla.
You can even dynamically define which skin is used by the application. To do this you would define that the skin name, in the trinidad-config.xml file, references a variable which could be set depending on some factor such as location, browser, logged in user or even let the user select their own skin!
Corporate identity is the persona of a business and many businesses invest heavily in the research, design and execution of that identity. Building that persona into your

Figure 7- Two applications

applications can be a critical element of application development. Oracle ADF has, at its core, the ability to customise the look and feel of an application through skinning. As Figure 7 shows, two applications built using the same technology can have wildly varying look and feel. This gives you, as a Fusion developer, the ability to easily deliver dynamic branding to your applications.

UKOUG AUTHOR PROFILE

Grant Ronald is a Senior Group Product Manager working for Oracle's Application Development Tools group responsible for Forms and JDeveloper where he has a focus on opening up the Java platform to Oracle's current install base. Grant joined Oracle in 1997, working in Oracle support, where he headed up the Forms/Reports/Discoverer team responsible for the support of the local Oracle Support Centres throughout Europe, Middle East and Africa. Prior to Oracle, Grant worked for 7 years in various development roles at EDS Defence.
Grant is author of the "Quick Start Guide to Oracle Fusion Development: JDeveloper and Oracle ADF", published by McGraw-Hill.

ORACLE DATA MINER COMES OF AGE

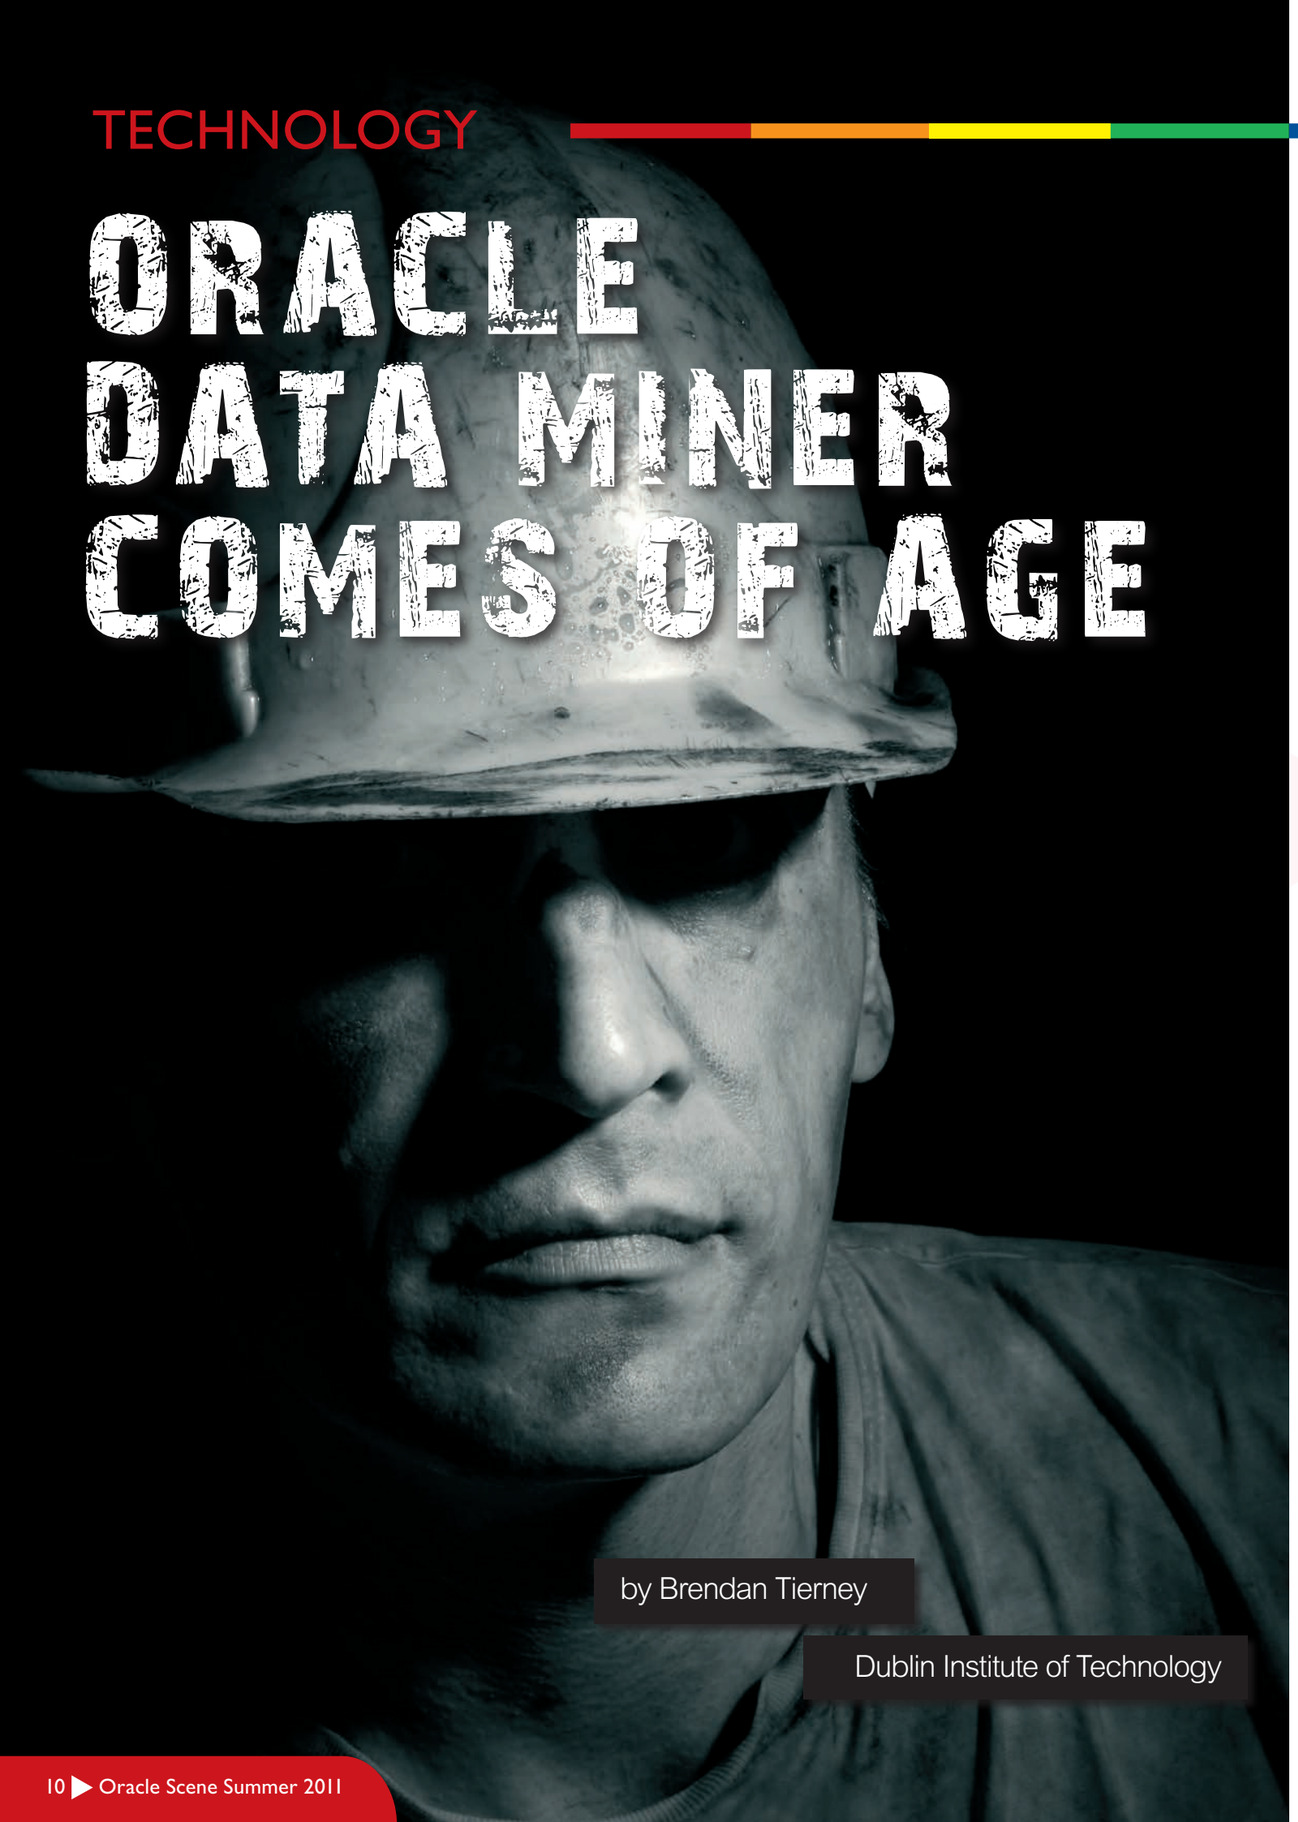

by Brendan Tierney

Dublin Institute of Technology

TECHNOLOGY

March of this year saw the release of Oracle SQL Developer 3. In this new release of the tool we have a significant number of new features, but more importantly we now have embedded, the previously stand alone products of Oracle Data Modeller and Oracle Data Miner. This article focuses on the Oracle Data Miner tool.

Oracle Data Miner (ODM) is now a work-flow based tool which has increased the functionality available to a data miner significantly, as well as improving ease of use. The new ODM tool is the first major release and upgrade of the tool in many years. In this article we take a look at the new features of the tool by working through an example of building a data mining model using one of the sample data sets provided.

Background

The Oracle Data Mining tool has been around since the early days of Oracle 10g, and I have been using the tool since 2005. The ODM tool up to now has been primarily a wizard type interface allowing you to analyse, build models, evaluate them and then applying these models to new data. The 10g version of the tool was simple in its approach and although it was very easy to pick up and learn, you had to perform a number of repeated tasks for each model you wanted to create. In addition to this, the outputs from the model build steps was very limited and it was a bit of a cumbersome task to try to evaluate the different models you had created.

The new ODM tool (11gR2 in SQL Developer 3) has addressed all of these short comings and the development team have really created a tool that is now at a comparative level with the likes of SAS Enterprise Miner, which is considered the number one product in the market. The ODM tool in SQL Developer is free to use, but you need to licence the data mining option as part of the Oracle 11.2g Enterprise Edition. If you take the list price for the Data Mining option in Oracle 11gEE, compared to the cost of purchasing SAS Enterprise Miner, plus all the added benefits of less data movement, in database data mining and the ease of deployment of the models, the likes of SAS and others are going to come under increasing pressure from Oracle Data Mining.

New Features

There are many new features in the new tool and these can be grouped under the following headings:

Data exploration: The first step of every data mining project involves investigating the data to try to learn from the data, gather some initial information and investigate if there are any patterns in the data

Workflow interface: This gives the user a more intuitive way to work with the tool and with the overall process of data mining. It allows for the repeated re-running of the data modelling process without having to input and define each step again. You had to do this in the previous version of the tool.

Generate multiple models at the same time: This is one of the major improvements in the tool. It allows you to create models using each of the algorithms available for each data mining techniques, in one step, instead of repeatedly defining each in the pervious version of the tool.

Graphical representation of models: Another major new feature. The tool now produces Decision Trees and Clusters graphically. With the Decisions Trees we can now see on the screen how the tree looks and then to investigate the different branches of it to see how the tree was built. We can also see what rules were generated to create these branches.

Evaluation of all the developed models: In the previous version of the tool you were presented with a set of evaluation diagrams and measures for each model. You were not able to see all the results on one graph and you had to resort to having multiple windows open at the same time to try to compare the results. Now we can get the evaluation measures and graphs for all the models on the one set of graphs. This allows a data miner to concentrate on determining the most appropriate model to use.

Each of these new features really deserve an article by themselves to illustrate their new capabilities.

Getting Started

Before you get started with using the new tool, you will need to install a data mining repository in your 11.2g database. There is an excellent step-by-step guide available on the Oracle By Example website for download called 'Installing and Setting up for Oracle Data Miner'. The new tool only works with the Oracle 11.2g Database Enterprise Edition. You have two options when installing the repository. The easiest is to allow SQL Developer to create the repository. But this requires you to have access to the SYS password. Depending on your environment you may not be in possession of this and you have to get your DBA to create the repository. The second option is to

run the individual scripts to create the repository. The two options for creating the repository will also set up some sample data for you to use while getting use to the tool. This sample data is based on the SH schema. We will be using one of the sample data sets created when setting up the repository in this article.

HINT : Make sure that the Samples were installed in the database.

The SH schema is used to set up the sample data for data mining demonstrations in the Oracle By Examples.

After you have created the repository the first step you need to perform, before you can start building models, is that you need to create a DM Project under your DM schema, and then a Work Flow.

Exploring the Data

The first step in any data mining project is to explore the data you are going to use. In the Component Data Workflow Palette you can select the Data Source option and drag this onto your work flow.

You can then select the table or view you want to use as a data sources and what subset of the columns you want to include. For this article we are going to use the MINING_DATA_BUILD_V view.

The next steps are to click and drag the Explore Data option from the Component Palette and then to connect the Data Source node to the Explore Data node in the work flow.

To run the work flow right click on the Explore Data node and select Run from the menu.

When the work flow has run and completed its steps, you can now explore the data and the gathered statistics. There are a range of statistics gathered for each attribute in the data set (Fig 1). In addition to the statistics a histogram is created for each attribute. From analysing

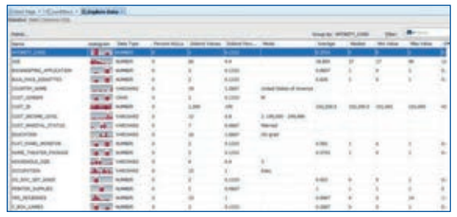

Figure 1- Explore data statistics

these statistics we can determine what attributes contain useful information and we can also start to work out how to handle missing data in each attribute.

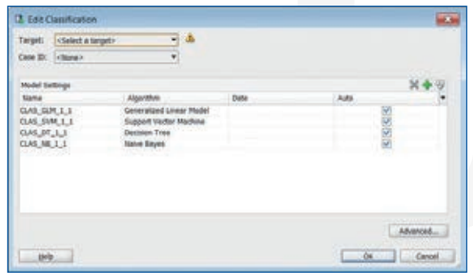

Figure 2 - Classification model section window

Creating a Model

To create a data mining model, the first step to determine what type of data mining we want to perform.

In our example we want to use Classification, but there are also Association, Anomaly Detection, Clustering, Regression, etc available.

To create a Classification Model node in our work flow, you will drag the Classification node from the Component Palette and connect it to the data source. There are four classification models available for you to run. You can select only one of them to run or you can select all of them (Fig 2).

You need to set the Target attribute. In our case it is AFFINITY_CARD.

Next you need to set the Case ID. This is typically the primary key attribute (CUST_ID). You can customise the internal parameters of each classification algorithm by selecting the algorithm and clicking on the Advanced button. When you are ready to run the Classification node, you right click on the Class Build node and select run.

Evaluating a Model

One of the major improvements in the ODM tool has been the ability to view the performance measures and graphs for all the models produced on the one graph or in the one window. To view the performance results from the different classification algorithms you can right-click on the Class Build node and select Compare Test Results (Fig 3).

The initial set of performance measures gives us the Predictive Confidence (%), Average Accuracy (%) and the Overall Accuracy (%) bar charts and corresponding percentage table for each of the classification algorithms. Other performance measure tabs available to us in this window include, Performance Matrix, ROC, Lift and Profit.

We can use all the performance measure information provided in these tabs to evaluate the models and to determine which one we will select as the model we want to apply to our new data.

One of the new features highlighted at the beginning of this article is that the ODM tool now gives

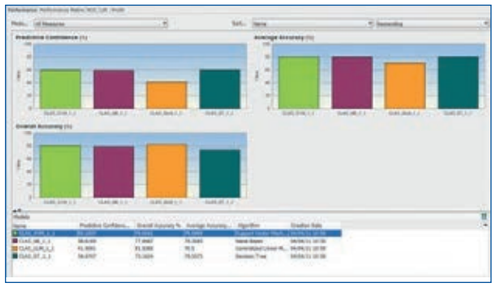

Figure 3 - Classification Performance Measures

a graphical representation of the output from the Decision Tree algorithm. To view the Decision Tree, you would right click on the Class Build node, in the work flow, select View Models from the menu and then select the decision tree

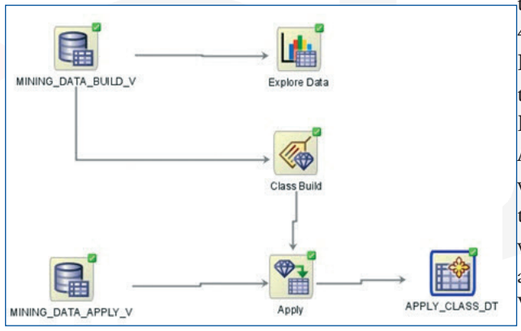

Figure 4 - Select Classification model to apply to the new data

model from the list. The Decision Tree will be opened in a new tab. We can now explore the tree, viewing the features of it like the different nodes and branches, the split rule, the number/percentages of slip cases and the underlying rule for each branch. If we selected the Decision Tree algorithm to deploy in our production environment, we can use the rules for each of these branches to implement in our code.

Applying a Model

After selecting which model gives the best performance we can now Apply this model to our new unseen data. The model will the work out, with a percentage of accuracy, what it thinks the class/target attribute would be for each record in the Apply data set.

To setup the Apply process on the work flow, the following steps need to be performed:

1. Create a Data Node for the Apply Data set;
2. Create the Apply node. You can click and drag the Apply node from the Component Palette onto the workflow;
3. Connect the Apply Data set to the Apply Node;
4. Identify the Primary key of the Apply data set. Double click on the Apply node. The window contains the two attributes that will be generated by applying the model. We also need to add the primary key column of the apply data set;
5. Connect the Class Build node to the Apply Node;
6. Select the Classification algorithm to run. Click on the Class Build node and view the details in the Property Inspector. Under the Output column, select/deselect the algorithms. In Fig 4, we select the Decision Tree model to run;
7. Create a data source to store the results. Click and drag the Create Table or View node onto the work flow;
8. Connect the Apply node to the New Table or View node;
9. Run the Apply process. Right click on the New Table or View node and select Run from the menu.

To view the output of applying the selected classification model to the new data, right click on the New Table or View node and select View Data from the menu.

Conclusion

The new ODM tool, which is part of SQL Developer 3, is certainly a major step forward compared to the pervious versions of the tool. With its new work-flow format, integration with SQL Developer, and new graphical outputs, amongst other new features, it is now a real competitor to the other major data mining tools, like SAS Enterprise Miner.

Given that there is less data movement required (compared to the other data mining tools) to get your data into the tool, combined with its new features and the features of SQL Developer all in one tool, Oracle Data Miner should now see a significant increase in market share.

UKOUG AUTHOR PROFILE

Brendan is a independent consultant and lecturer with the Dublin Institute of Technology in Ireland. Brendan has extensive experience working in the areas of Data Warehousing, Data Mining, Data Architecture and has worked on projects in Ireland, UK, Belgium and USA. He has been working with the Oracle database and tools since 1992, starting with Oracle 5, Forms 2.3, ReportWriter 1.1.

email : brendan.tierney@dit.ie

twitter : www.comp.dit.ie/btierney

www.comp.dit.ie/btierney/DataMining/Notes

The Joel Goodman interview

In Conversation with Oracle Expert Joel Goodman

Technical Team Leader, Oracle Core Delivery Europe, Middle East, Africa.

OS: *Joel, let's start by explaining your role within Oracle?*

Joel: Core Delivery is the part of Oracle University that delivers courses on the core technology stack - from the middle tier application server to the database, storage and operating systems layer. It doesn't include applications. It includes languages and tools, and everything that uses those in the Middle Tier.
I don't have expertise in all those areas but, from a technical management point of view, I'm the Technical Team Leader advising management on technical strategy. I spend about a quarter of my time doing seminars and courses for Oracle University. About a third of my time is spent working with the Oracle Certification Team, leading the development of certification examinations for Oracle Database Administrators. This includes Oracle Certified Associate, Professional, Expert and Master Exams.
I also work fairly closely with the Curriculum Group on helping to design and develop the database and Linux systems course content, specialising in Dataguard, RAC, Grid Infrastructure, and in the past two years, the Oracle Exadata Database Machine.

Oracle Expert - Joel Goodman

OS: *Could you give a brief overview of Grid Infrastructure installations, for example, for clusters or standalone servers?*

Joel: Sure. Firstly, there are two flavours of Grid Infrastructure; for a cluster and for a standalone server. In order to understand this we should go back to the origins of this software component.
Grid Infrastructure consists of two separate components in each of those flavours. The origins of Grid Infrastructure start with the Clusterware from previous releases. Clusterware belongs to a family of products known as high availability software (HA) which Oracle has provided since Oracle 10g Release 1. Other operating systems vendors have their versions of HA software. For example, Sun cluster for Sun and HACMP for IBM AIX and HP Serviceguard. Clusterware allows managing multiple nodes as a cluster and also provides for failure isolation of I/O from failing nodes. It also provides tools for the management of resources that run on particular nodes in the cluster - including re-starting place or failing over a resource to surviving node is required in the event of node failure.
Oracle Clusterware has been a pre-requisite bit of software since Oracle 10g, to run RAC databases. Traditionally the OS Administrator managed High Availability software, although Oracle Clusterware is sometimes administered by DBAs.
When Oracle Clusterware was originally released, it had dependencies on other technologies.
One of the requirements for running a cluster is some form of shared storage. You can have shared storage provided

TECHNOLOGY

in using Network Attached Storage (NAS) on filers. There is also the possibility of using shared storage from Storage Area Networks or SANS and traditionally on RAC databases, the Clusterware shared storage was provided by one of these methods, usually by a SAN.
Oracle 10g Release 1 also contained Oracle's own storage technology called ASM - a shared storage solution for Oracle databases.
The problem was that the Clusterware itself has certain files that are required for it to operate. It has the OCR - the Clusterware repository files - which contain a little repository describing the Clusterware, the nodes of the cluster, the resources that run on the cluster, which nodes they're meant to run and so on. These must be shared files but couldn't be stored in ASM. There are also voting files which need shared storage but these couldn't use ASM in Oracle 10g or in Oracle 11g Release 1. This meant that customers running Clusterware required some other shared storage solution such as using SAN or using filers in addition to using ASM if they were using ASM for their database shared storage. One of the design requirements for 11g Release 2 was to have ASM become a shared storage solution for every possible requirement including the Clusterware. To do this required that ASM be available during the installation of the Clusterware to allow the Clusterware files to be defined as residing inside ASM. That required ASM to be installed at the same time as the Clusterware was being installed otherwise you wouldn't have ASM there to begin with! The consequence of this was that the installation of the Clusterware would have to include the installation of all the executables and other files needed to support ASM as well. So, this couldn't be called this 'Clusterware' anymore, because it combined technology from two different areas - one was the ASM technology which originated as a database oriented technology and one was the Clusterware which originated as an operating system administrator orientated technology. A new name was needed to describe this mix, and the result was Grid Infrastructure. It provides the infrastructure for grid computing including RAC, which is the database layer for grid computing.
That explains why there was a name change and it also explains one of the very important new features in Oracle 11g Release 2 which is that ASM can be used not only for RAC database storage but shared cluster storage as well. And, that's not the only shared storage requirement. There are others, which I'll get to when we discuss some of the ASM features of 11g Release 2.
Inside the Clusterware there's a component called the Cluster Ready Services daemon, which is the bit that manages this and its job is to start, stop and monitor the resources that are meant to run on the various nodes of the cluster.
It can make decisions about restarting failed components in place, failing these resources over to other nodes and maintaining metadata dependency relationships. A resource which depends on another resource to be active for it to work properly, will have that resource automatically started and the Clusterware will not start something before one of its dependent resources has started as well.
In Oracle 10g and 11g Release 1, if a customer wished to use ASM, they would install another Oracle Home from which they would typically run their ASM instances. ASM, in a non-clustered environment, still requires a very small subset of Clusterware to do specific work involving resource locks and a few other types of cashing mechanisms so that the ASM environment can coordinate its activities with the database instances running on the same machine. This is done for releases prior to 11g Release 2, by having a small subset of Clusterware running called Group Services even when running on a standalone server. Grid Infrastructure for a standalone server, still includes this functionality, but also contains the resource management functionality of the Clusterware with no failover capability. But it can automatically start resources in the correct order, monitor their health and restart them in place; all nice extras for a standalone server that can help customers to avoid having to do their own implementations.
Even though there are no other nodes on a standalone server, there's still the possibility of:
A) Having resources start up automatically;
B) Having them shut down cleanly;
C) Monitoring them and then restarting them in place as appropriate.
Prior to 11g Release 2 DBAs would need to script this themselves. There were some tools from Oracle to automate start-up and shut down of databases but there was no monitoring capability built in. Oracle Restart is the name of the Clusterware subset that's used for managing resources on a standalone server. It not only manages start-up in the correct way but provides a comprehensive set of administrative tools. Of course, many customers work in both environments - a combination of standalone Oracle servers as well as clusters. Now that Oracle Restart exists, certain new features that have been developed in 11g Release 2 require that the customers use Oracle Restart. I predict that, as time goes by, more and more features will depend on customers having Oracle Restart installed because it's there now and if it's there now it's there to be used. So it's an ease of use solution, it's a way of making administration consistent and it's very much the case that most people can benefit strongly from the automated management and monitoring rather than having to write their own scripts to provide the solution.

OS: ***Given that the install process and the set-up of Grid Infrastructure is well documented, what are some good tips and pointers would you say, for troubleshooting cluster problems?***

Joel: It's all about preparation. Installing Grid Infrastructure is more complex than a Database Administrator might expect. Grid Infrastructure has many pre-requisite steps to make certain that the cluster is in a proper condition for the Grid Infrastructure install to be successful. It's more the sort of things an Operating System Administrator would be familiar with, but then they wouldn't know some of the Oracle specific requirements. We spend quite a bit of time on preparation and configuration during our Grid Infrastructure administration classes, to help customers appreciate the need for planning.
A hefty chunk of time is spent doing installs of the Grid Infrastructure, so students learn not only to do the install, but to recognise the kind of set-up that's required. Fortunately, Oracle provides two tools that might help administrators handle this. We have a Linux package - an RPM- called Oracle_Validated. Installing this will make the operating system 'Oracle ready' to some degree meaning it will create certain users and set certain OS kernel parameters to appropriate values for an Oracle install.
It is really designed primarily for single instance installs not for RAC but can still save some time when installing Grid Infrastructure for a standalone server.
For clusters, there is a tool called the Cluster Verification Utility (Cluvfy). It can be used in a variety of contexts. It can be used at different stages of the proceedings to check whether or not the environment is correct. It can also be used, after everything has been installed successfully, to check different components

> "Core Delivery is the part of Oracle University that delivers courses on the core technology stack - from the middle tier application server to the database, storage and operating systems layer."

TECHNOLOGY

of the Clusterware. There are stage checks and component checks in the tool. So, from an installation point of view, my advice is to always follow the requirements in the documentation.

Installation will:

- Require a certain amount of shared storage;
- Require network adaptors such as a public network adapter for each node;
- Require another network adapter for an interconnect between the nodes;
- Check for certain OS users;
- Check for the presence of certain packages on the operating system; and,
- Conduct a post HWOS stage check (post hardware and operating system) to check if you are ready to start installing your Grid Infrastructure.

There are other stages along the way, which students cover in class, but this should be your first approach - be prepared and think it through. The Grid Infrastructure installation for a cluster has certain optional components which may be installed using the advanced install which have other set-up requirements. But I think it's beyond the scope of this little discussion to do that in great deal of detail. One of the features, called Grid Naming Service (GNS) is in effect, a Domain Name Server and knowledge of the administration of DNS, servers, DHCP servers and networks is required to use some of these extra advanced features.

This makes GNS in 11g R2 very interesting but it also creates certain challenges because the scope of the technology stack is quite extensive. It is doing some things which are OS Administrator oriented. It's interfacing with Network Admin tools. It's also doing other things involving shared storage which might involve Storage Administrators. It requires a broad multidisciplinary skill set to the point where, at the very least, the people responsible for the install can hold an intelligent discussion with their colleagues.

OS: ***Is GI functionality better than the proceeding Oracle Clusterware and automatic storage management functionality?***

Joel: You can run Clusterware without using ASM if you don't want to. You can use filers for example. The new Oracle 11g Release 2 Clusterware has several new features. One of them is a new approach to managing the Clusterware itself. In previous versions, there were certain platform dependent implementations for getting the Clusterware stack initialised when the operating system itself was booted. So, if you were running on Linux then it might be done in a certain way. If you were running on Solaris it would be slightly different and so on. Oracle made a decision that it was easier to have a consistent way of managing the start-up, shut down and monitoring of the Grid Infrastructure components themselves, so what they've done is create a new layer called the High Availability Services Demon (OHASD). OHASD's responsibility is to get the Grid Infrastructure stack started up on a node. OHASD is generic so when an operating system reboots, OHASD is started, it looks to see whether the Clusterware is meant to run enabled or disabled. If the administrator runs it disabled, it means the Clusterware is only started if you start it manually. If its running enabled then OHASD will start it automatically. But either way once you get the Clusterware started OHASD starts up the components of the Clusterware in the correct sequence using the same kind of logic used by CRSD to manage the resources that are themselves managed by the Clusterware.

CRSD manages resources that can be cluster-wide resources, meaning resources that can either run on multiple nodes or resources that may fail over from one node to another.

Interestingly, when you run Grid Infrastructure for a standalone server there is no cluster, there is only the local node and therefore OHASD does everything. There is no CRSD when you use a standalone server because there's no cluster. The code for OHASD and the code for CRSD are almost identical, since they're more or less doing the same kinds of things; they're just doing them to different types of resources. In previous releases, for example, if I wanted to check the state of the Clusterware itself on an 8 node cluster, I'd have to connect to each node and run the crsctl command. Now I have an OHASD daemon that talks to its counterparts on other nodes. I like that feature particularly but there are others. The number of different resources managed by the Clusterware within the Grid Infrastructure, for example, is far more granular than it used to be. For example, in older releases, we knew that database instances existed and we knew that they could depend on the ASM instance located on the same node. If the ASM instance was not running, then the database instance wouldn't be started because it would not be able to access its data. However, if the ASM instance was started, but one or more dependent disk groups were not mounted, then errors would occur because the granularity of dependency was at the level of disk groups and not simply to the ASM instance. ASM disk groups are equivalent in some ways to a logical volume and if you have an ASM that's running and there are five disk groups but three of them are mounted and two of them are not mounted then any database instance that depends on one of those disk groups that wasn't mounted could get an error if it tried to start up. The trouble was that, in previous releases, the granularity of resource management was restricted to the ASM instance as a whole and not to the individual disk groups that it managed. That's now been changed in Oracle 11g Release 2. Each disk group is a separate resource and the dependency relationships that are defined in the OCR are down at a level of the disk groups upon which a database instance depends.

> "CRSD manages resources that can be cluster-wide resources, meaning resources that can either run on multiple nodes or resources that may fail over from one node to another."

When the Clusterware wants to start a database instance on a particular node, it needs to ensure the ASM instance is up on that node and that specific disk groups, needed by the database instance are mounted. So the granularity of resource dependency is another change in the Clusterware that I quite like.

Another new feature that I like is called the Single Client Access Name (SCAN). The SCANs and SCAN listeners have made the architecture of the Oracle listeners very different. They provide a separation of responsibility for listeners between those listeners which handle connection requests and those which make the load balancing decisions. It presents the possibility that the two different listener types will end up diverging based on their differing responsibilities. Right now they're still running the same executable.

One of the new features for ASM that I would point to is the ability to support the voting files and the OCR in ASM. ASM now provides a shared storage solution for every conceivable type of shared storage needed on an Oracle cluster.

The Grid Infrastructure also supports the ability to manage third party non-Oracle applications such as Apache based web applications. They may be monitored using the Oracle Clusterware and are provided with the same kind of high availability and fail-over as are provided for Oracle RAC database instances, listeners etc. These applications may also require shared storage on conventional file systems. If a web application runs on node 1 of a cluster and it accesses data, but then fails over to another node it still requires access to the same data as it was using on the original node. ASM did not support that in previous releases.

Oracle database homes are stored in a flat file system too, but older versions or Oracle did not provide general purpose cluster file systems stored inside ASM. Grid Infrastructure permits this, allowing an Oracle database home to be stored inside ASM. This can help when single instance Oracle databases are supported under the Clusterware for cold failover, and the home directory

is only mounted on one node at any one time. The file system is stored inside ASM, and is also defined as a cluster resource on which the database instance depends. This guarantees that the file system is mounted by the Clusterware before attempting to use any files in the home directory.
ASM was extended in 11g Release 2 to include the ASM Cluster File System (ACFS). The space for files in the file system is allocated from ASM disk groups rather than being provided from device files or luns. You can use ACFS for a shared Oracle home directory; you can use them for application files, or for DBA oriented files such as the directory specified in the diagnostic_dest parameter. ASM now provides a complete shared storage solution so that it's possible to run the Grid Infrastructure without using any other shared storage solution. There are some other very nice features in ACFS as well. ACFS was available in Oracle Release 11.2.0.1 for Linux and Windows only. From Oracle Release 11.2.0.2, it's also available on Solaris. It supports copy on write (COW) snapshots, something that traditional Linux EXT 3 file systems didn't have and the Windows file system didn't have. Furthermore in Oracle Release 11.2.0.2, ACFS has been extended to support replication as well where an ACFS file system on one node is replicated automatically to another nodes It has its own logs and its own method of transport to the remote location providing a protection mechanism in case of failure at the original location.

OS: ***What are the Oracle University courses covering this kind of content that you'd recommend?***

Joel: There is a 4 day course on Grid Infrastructure Administration that covers the install of the Grid Infrastructure for a cluster. It includes monitoring, administration of the Clusterware and administration of the ASM as plus a bit of diagnostics and tuning. The course is aimed at OS Administrators, Storage Administrators and Database Administrators. In my experience, the attendees have been primarily OS Administrators with a few DBAs attending too. That skill set is a prerequisite for the three day RAC Course - a stand alone three day 11g Release 2 RAC Administrators Course.
To administer RAC databases assumes some knowledge of Grid Infrastructure. So if people have already got RAC administrators skills, then they only require Grid Infrastructure administration skills. They should attend the 4 day Grid Infrastructure course only and not bother with the 3 day, RAC admin course, because 11g Release 2 RAC isn't very much different than 11g Release 1. There are one or two changes involving services that really make a difference. Someone with minimum or no RAC skills would probably need to attend both courses.
Oracle has another option, but it's challenging. There's a course called the Grid Infrastructure and RAC Administration accelerated course which is a five day event packing in seven days of content. The first day tends to run till 6:30 or 7 o'clock at night. The second day is also long but and in the first two days we probably cover nearly three days worth of work by having an extra four or five hours. I also assign a bit of homework in the middle of the week. The 'boot camp' is a popular choice. In the Americas where travel distances are considerable, they don't run the four and three day course separately very often if at all; they tend to run only the accelerated version.

Joel Goodman is Oracle's Senior Principal Technical Team Leader.

For more information about Oracle University contact:
T: 0845 777 7 711
T:+44 11 89 726 500
E:edenrollment_uk@oracle.com

Application Server & Middleware

by Simon Haslam

March saw an innovation for the UKOUG Application Server and Middleware SIG: a new "dual-venue" format. Working to a single agenda we had some speakers, at each location, with each presentation being relayed live to the other site.

The venues were Oracle's City Office in London and Fujitsu's office in Warrington (thanks to both for providing such excellent facilities) and our theme for the day was "Oracle Fusion Middleware 11g Upgrade" – a subject close to many people's hearts in light of the end of 10g Premier Support later this year. The plan was to transmit the slides via a live webcast and broadcast the speech over a conference phone line (not wanting to put all eggs in one internet basket!). We also distributed the slide decks to both sites in advance so our fallback would have been "next slide please," or perhaps "beeeep" like the audio-visual presentations in French lessons (for those educated in the 1980's!).
On the day, the technology worked almost perfectly, with overhead loudspeakers in Oracle's office meaning that even people at the back of the room could hear clearly. Pushing the technology we attempted a networking session, though without roving microphones this wasn't quite so successful – I've a few ideas up my sleeve for the next time!
One of the reasons many people had attended the SIG was to hear John Stegeman's Forms 11g Upgrade presentation. Unfortunately family circumstances, the previous day, meant that John wasn't able to appear in person but, thanks to the technology, he still delivered his presentation, much to the appreciation of the delegates.

The post-event feedback was generally positive. Only one delegate (in Warrington) felt the new format was "much worse" than a normal, single-venue SIG; 23% felt it was "slightly worse"; and the rest thought it was the same or better. In addition nearly 60% said that the dual SIG format was "slightly better" or "much better" than webcasts delivered to people's desktops. Notably, feedback from London (where 5 of the 8 sessions were delivered from) was generally more positive than Warrington, suggesting that we need to improve the experience at the secondary site.

A few other comments from people there on the day: "all the sessions were geared perfectly around the issues we have", "we wouldn't get 1:1 or 1: few discussions at coffee/lunch and subsequent swapping of email addresses with people if done as a webinar", "the SIG format still has the networking advantage", and "I was surprised and pleased at how well this worked".

Thank you to all of those who took the time to complete the survey.
Given that we made the SIG much easier for some members to reach, and increased our attendance numbers by more than 25 per cent (which I suspect could be higher for other locations), we considered this a successful experiment. By the time you read this, we will have had our second dual-venue SIG (Birmingham with London as a secondary, and Virtualisation as its theme). Hopefully we will have been able to fine-tune the format and the event will have had a bumper turnout!
This is just one of the ways the SIG committees are trying to make the events more accessible and valuable to members. If you have other suggestions please send them to: ideas@ukoug.org

UKOUG AUTHOR PROFILE

Simon Haslam is in his 4th year as Chair of the UKOUG Application Server and Middleware SIG, and is a regular speaker at Oracle events. He's one of only two Oracle ACE Directors (Middleware & SOA) in the UK and an enthusiast of all things to do with Oracle infrastructure. His time is split between middleware consultancy and troubleshooting, so he can be usually found in his native habitat of 'the customer meeting room' (scrawling on a whiteboard) or else frantically tapping at a keyboard! Simon's Fusion Middleware blog can be found via: http://simonhaslam.co.uk/

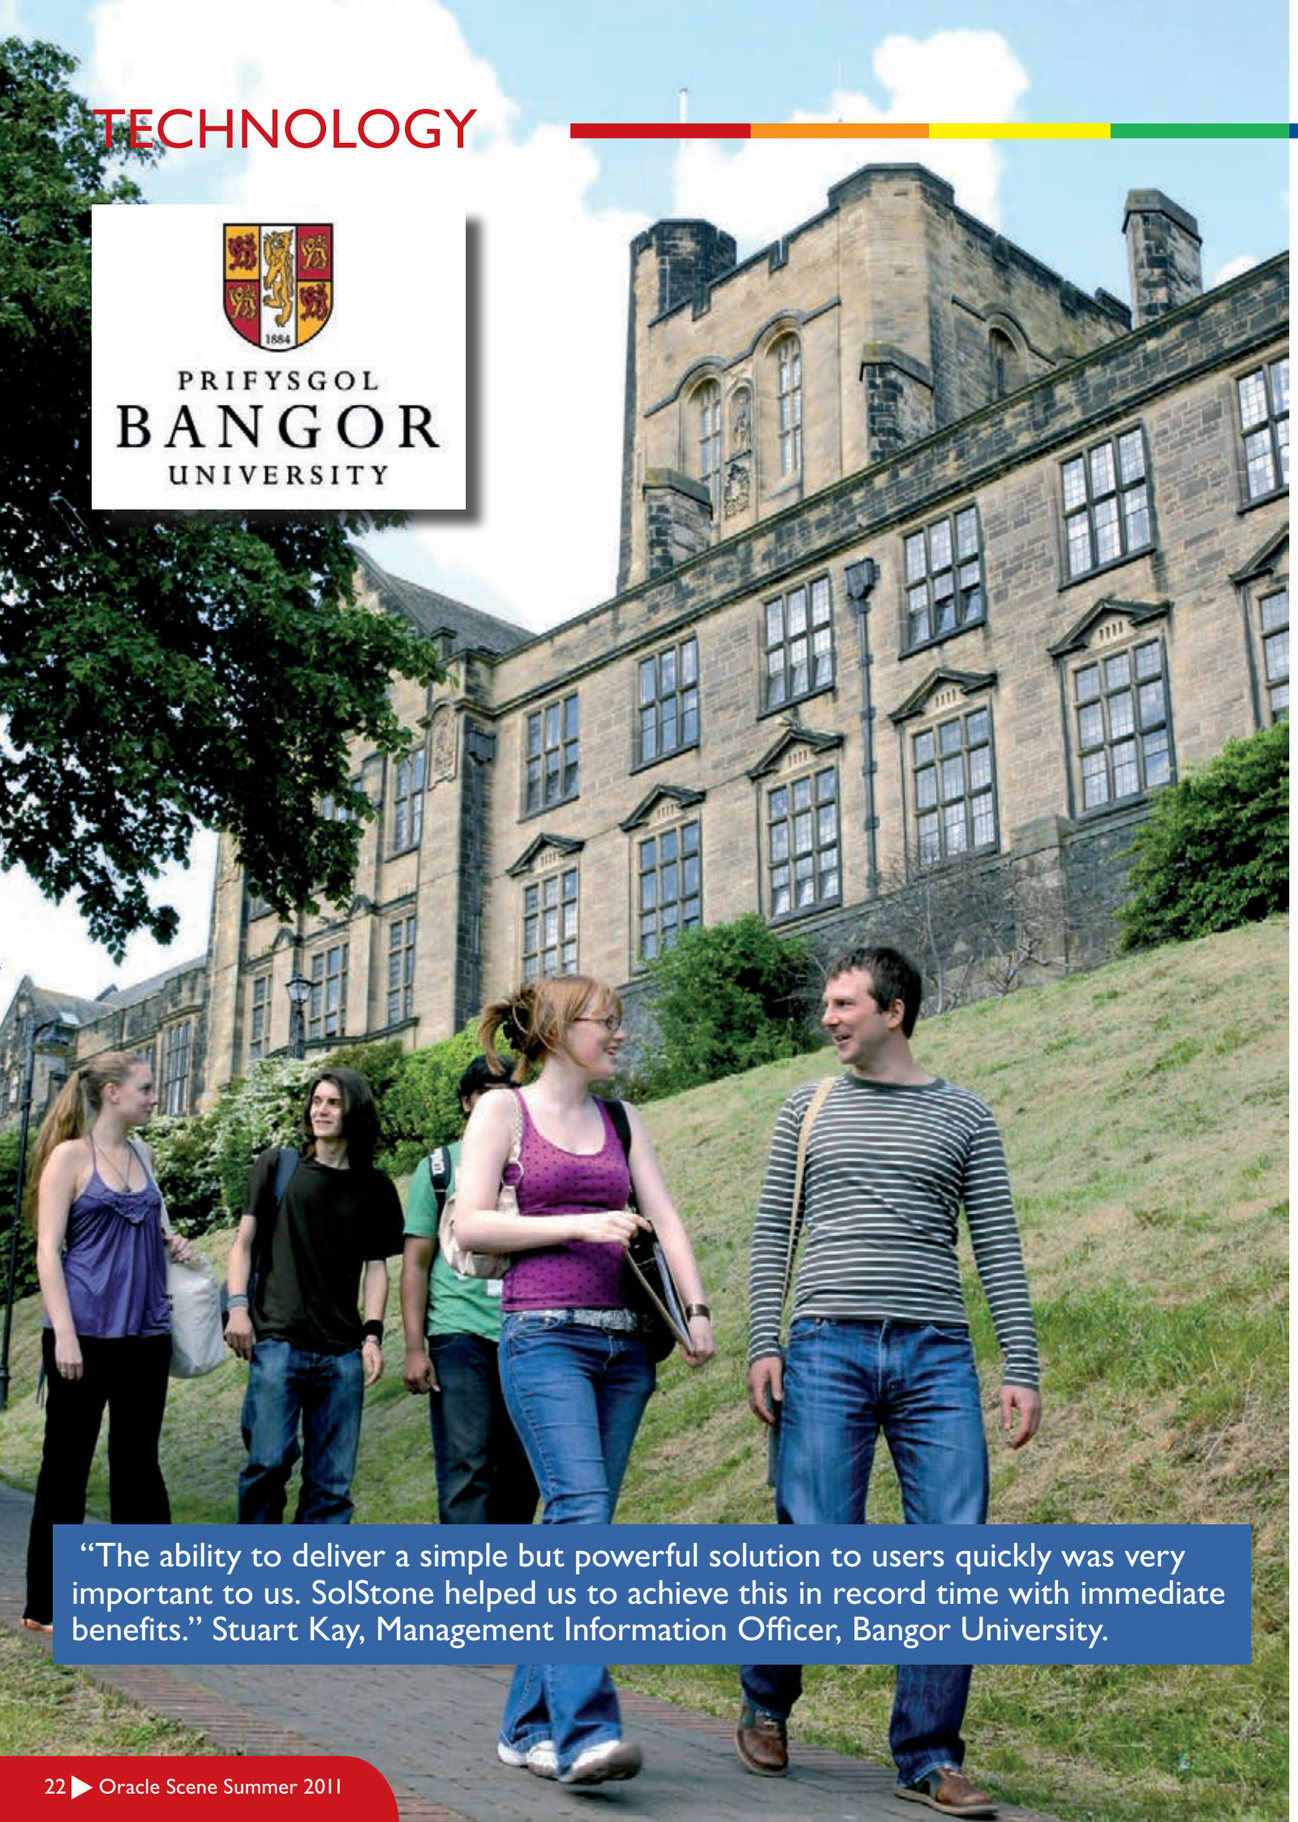

advertorial

Bangor University Achieves Institutional Intelligence from Oracle BI

With increasing governances and fiscal demands, Bangor University needed a way for key staff to quickly and easily track critical student, staff and financial data. After implementing Oracle Business Intelligence Enterprise Edition (OBIEE), stakeholders now have access to near real-time data to support effective decision making.

Thanks to their new Oracle Business Intelligence (BI) solution, instead of spending the majority of their time locating and gathering data, stakeholders can now deal with critical projects that add more value to the Colleges, Schools and Departments of the University.

Background

Established in 1884, Bangor University has a distinguished history and significant record of achievement. Bangor University currently has 12,000 students and an annual turnover of £126m with twenty three Academic Schools arranged in five Colleges. As a leading research-led University with an international reputation for teaching and research, the University is recognised as a centre of excellence for a varied portfolio of academic programmes and for the high quality experience it provides for its staff and students.

In order to keep the University running effectively and efficiently, it employs a number of best of breed systems, including Banner, Agresso HR and Agresso Finance.

With disparate systems, it was difficult to see the big picture performance of the University or have the ability to drill-down to see necessary student, staff, financial and research data. In addition, creating comprehensive, timely reports on aggregate student and staff data across a range of variables and timescales was also a tremendous challenge.

The University needed to find a simple and timely way to gain accurate, detailed statistics on student recruitment rates and profile students according to age, origin, ethnicity, and course in order to determine the best mix of study programs.

Another challenge for the University was to find a way to monitor trends in student and staff data across the Colleges, Schools and Departments of the University.

During the tender phase, Bangor University was introduced to the Oracle BI solution by Oracle Platinum Partner, SolStonePlus. Chosen for its successful track record in delivering performance management solutions in the higher education sector, SolStonePlus was tasked with designing and developing a BI platform that would allow all key staff to instantly access the information they needed - a platform that could scale to support the University's growing needs with

the capability to support potentially hundreds of users.

The Solution

To the University, the Oracle BI Suite was the most cost-effective, fully integrated and easy to use solution on the market.

Together, the teams quickly began to build the solution, leveraging the Oracle Database 11g and Oracle Warehouse Builder 10g. The system was tailored to meet the specific needs of Bangor University and designed in such a way that maintenance of the metadata is simple and the full power of the reporting tool is leveraged with ease.

In approximately 70 days, SolStonePlus was able to design and develop the initial data loads, train the development team and release the first version of the Oracle BI solution to 50 academic and managerial decision-makers within the institution.

"The ability to deliver a simple but powerful solution to users quickly was very important to us. SolStone helped us to achieve this in record time with immediate benefits." Stuart Kay, Management Information Officer, Bangor University

"SolStonePlus has enabled us to build a pervasive real-time intelligence environment for monitoring student recruitment, student and staff activity that has the scalability to support enterprise-wide reporting as our needs evolve."
Stuart Kay, Management Information Officer, Bangor University

The Benefits

With the latest intelligence data provided to decision-makers within 24 hours instead of up to 4 days using the existing legacy tools, the University now has a system that immediately alerts them to potential business issues and performance shortfalls. It will provide a single, consolidated view of all their disparate management data, presenting the most relevant information in an easy-to-understand graphical format, enabling them to quickly drill-down to the root cause of specific business issues.

Furthermore, the system allows end users to select and build personalised data maps in order to build reports relevant to their areas of responsibility.

The University now has harmonised reporting of recruitment, student and staff data, made available via daily dashboards in both English and Welsh languages. The multilingual capability of OBIEE proved incredibly useful for Bangor University as an organisation operating in a bilingual environment, the system offers the ability to add in translations into the system.

Improvements in data accuracy were another added benefit of the system that allowed data quality issues to be highlighted and prioritised. Feedback from staff has been excellent and staff across all Schools and Colleges now have the opportunity to become focused and harmonised having true confidence in the data, empowering users at all levels to make more effective, informed decisions.

Next Steps

The University is extremely enthusiastic about the new system and is keen to explore its flexibility and leverage its full potential.

The team is currently working on a range of exciting new projects that are only possible thanks to the increased functionality of their new BI platform. These include using the solution to:

- Plug into the Student Destination Survey, in order to monitor employability and trends;
- Leverage near real-time intelligence during Confirmation and Clearing to profile applicants, confirm offers faster, and keep acceptance levels in prescribed range of government student number controls;
- Provide projected student number updates regularly during the day throughout the clearing period; and,
- Integrate student, staff and financial data to monitor research activity and quality.

"The good working relationship with SolStone has meant that we have been able to deliver Bangor specific solutions whilst building on their knowledge and experience. I have every confidence the project will deliver its goals and will provide the intelligence required by the University to deliver and monitor its strategic objectives," Dr Kevin Mundy, Director of Planning, Bangor University.

About SolStonePlus

The decision to implement any software solution or platform should never be taken lightly and will require the expertise of an experienced consultancy to ensure that your investment is leveraged effectively. SolStonePlus' highly experienced team can provide the support required to ensure you have full confidence in your investment.

With over 28 years experience across a wide range of sectors and solutions, SolStonePlus is an industry leader in the fields of Business Intelligence, Data Warehousing and Enterprise Performance Management solutions.

SolStonePlus has built an unrivalled reputation through the delivery of a broad range of high quality solutions that enable their customers to achieve vast improvements to their business performance including more proactive decision-making through access to higher quality, readily available data as well as better long-term planning through budgeting, forecasting and timely, accurate reporting solutions.

ORACLE Platinum Partner
Specialized
Oracle Business Intelligence Foundation

TECHNOLOGY

The GlassFish Tale

MARKUS EISELE

The former Sun flagship product GlassFish got its last major version update - number 3 - with the publication of the Java EE 6 release back in December 2009. Since the acquisition by Oracle, a lot of rumours have been circulating about what will happen to GlassFish. Amongst others, priority concerns were whether GlassFish would get its former clustering capabilities back, which had been included in the 2.x releases. The now-released version 3.1 marks a further milestone on the GlassFish community road map which was presented by Oracle early last year.

The community was speaking with one voice at the beginning of March 2010: Oracle will degrade GlassFish to a simple reference implementation (RI) and make it a children's edition! Due to its lack of cluster support GlassFish will no longer be suitable for productive use. The last point in particular was driving more and more customers towards GlassFish as a real alternative to commercial products such as WebLogic Server.

Today's 3.1 release ends speculation about the future of the popular open source server, as Oracle continues to fulfil the promises it made back in early 2010. The 3.1 project started on 19th of April 2010. After about 10 months, 8 milestones and a total of 43 promoted builds, as well as thousands of bug reports and fixes, a modular GlassFish Java EE 6 kernel was built, which provides clustering, centralised administration and high availability. It combines the advantages of 3.0 and 2.x and enriched them with many new features. It is based, like its predecessor, on the HK2 (Hundred Kilobyte Kernel) microkernel and therefore provides an easily extensible, flexible platform.

Clustering and High Availability

The most requested feature from the past year was clustering. At first glance nothing obvious seems to have changed in comparison with the latest clusterable 2.1 release but there are some major improvements waiting to be discovered.

One of the most obvious changes is the elimination of the so-called "nodeagent", which was the central hub for the cluster communication in 2.1. 3.1 introduces the centralised administration concept and uses two different ways to communicate within the cluster. The synchronization of server configurations and applications happens directly via Secure Shell (SSH). The Domain Admin Server (DAS) can connect via SSH to remote machines and either install or synchronise files. If SSH is not an option it's still possible to make the required changes directly on the remote nodes. All runtime cluster notifications are managed by the Group Management Service (GMS). This component, which is based on the Shoal Framework, notifies the DAS and is responsible for the individual state of instances. It also takes care of other basic cluster functions (eg, in-memory session replication, transaction services and timer services). In direct comparison with the GlassFish 2.1.1 cluster performance was raised by 34% [1]. This is also true for the maximum number of possible instances in a single cluster.

Developer Productivity

GlassFish provides many improvements for developers. Start-up and deployment times were reduced on average by 30%. This is achieved most dramatically by the dynamic loading of modules in the HK2 kernel. With a total number of about 262 individual modules, the cumulative speed increase is an achievement. The 3.1 version is able to restore the HttpSession State after redeployment. This feature has been popular among developers of complex application functionality for some time. It has now reached Stateful Session Beans and timers. It makes troubleshooting complex error situations much easier. Even in the integration with the development and build environments, things have moved forward. The embedded GlassFish has a new API and the associated Maven plugins have been updated to simplify automated in container testing. Eclipse and NetBeans are still the major development environments which are supported first hand. The freshly released NetBeans 7.0 has strong support and already has the latest GlassFish 3.1-b43 bundled. For Eclipse the server adapters have been extended and an even tighter integration is possible with the Oracle Enterprise Pack for Eclipse (OEPE).

Management and Monitoring

The topic of management and monitoring has been extended to the entire cluster. One of the main features of the centralised administration can be utilised by the use of the newly introduced RESTful API which allows the remote control of complete domains. As the name implies this feature exposes administrative and monitoring resources to HTTP clients with responses served in XML, HTML or JSON formats. Pretty much anything you can do with the web console or asadmin, you can do with this. The JDBC monitoring has also been expanded.

In addition to the statement leak detection, you can now trace SQL queries extensively. Connection pool usage can be monitored on an application level. Connection pool validation can happen based on custom validation templates. Beside the monitoring improvements, you'll also find some more management advantages. First and finest is the new "Application Versioning" system. It allows multiple versions of the same application to exist in a GlassFish domain. It provides the developer and the administrator extensions to the relevant existing tools and commands to deploy,

TECHNOLOGY

view, and manage multiple versions of applications. At most, one version of an application may be enabled on any target. Another newly introduced feature is the possibility to have so-called Application scoped resources. For this purpose, its own deployment descriptors have been provided, in which the configuration of required server resources (e.g. JDBC, JMS) can be placed. The resources will be registered during application deployment and unregistered on undeployment. Nearly mandatory are security updates and enhancements. Unix-based operating systems can now utilise the PAM (Pluggable Authentication Module) realm, which maps users from the operating system. For applications with extended client authentication needs, the new extended CertificateRealm has been added, which allows specifying an optional JAAS LoginModule to perform authentication. In addition to those two prominent added features, a large number of small security improvements have also been incorporated [3]. The discussion about the purpose of the coexistence of both application servers (GlassFish and WebLogic) is still ongoing. Oracle's announcements from the past seem to indicate that GlassFish is positioned as the Open Source Java EE Server while WebLogic remains the complete corporate solution with the tightest integration into the complete Oracle Fusion Middleware Stack. The 3.1 GlassFish release is the first one to position itself as a new development base for WebLogic. It is able to read and interpret one WebLogic specific deployment descriptor, the weblogic.xml. It seems as if there are more to come. And even WebLogic will be able to read and understand GlassFish specific configuration files very soon.

> "The 3.1 GlassFish release is the first one to position itself as a new development base for WebLogic."

Technology Updates

Away from the innovations, several Java EE 6 technologies have also been updated. JSF has been increased from 2.0 to 2.1. Also, the Contexts and Dependency Injection (CDI) RI called Weld has been updated to version 1.1.0. JPA RI EclipseLink is now included with version 2.1 and the underlying OSGi runtime of GlassFish (Apache Felix) has been brought up to a newer version (3.0). Only with the Open Source Edition is the first support for hybrid (OSGi enabled Java EE) applications based on various Enterprise OSGi Specs (RFC 66: OSGi/Web container, RFC 98: OSGi/JTA, RFC 122: OSGi/JDBC, RFC 142: OSGi/JNDI, RFC 143: OSGi/JPA) available.

Last but not least, Grizzly and Jersey have been raised to new versions, which contain a lot of bug fixes. In particular, the Grizzly update brings the long-awaited WebSocket support into GlassFish's webstack. One last area of technology updates is very internal. More or less silently, the HK2 improvements have been laid as the common platform for both Oracle application servers. It introduces Startup services, service trackers and service listeners together with a JUnit integration. The main work involved accommodating WebLogic Server requirements for adopting HK2.

Commercial Extensions

In addition to the open source bits, there is a commercial version available. It includes some closed source add-ons. They are bundled as GlassFish Server Control which includes functions for backup and recovery of the DAS. The Performance Tuner analyzes the underlying infrastructure and optimises the runtime settings of the GlassFish instance to achieve better throughput and scalability. The monitoring client scripting interface is the new solution for ad-hoc monitoring. It executes JavaScript programs against more than 100 fine grained server probes. For Solaris-based installations in combination with the JDK 7, an integration of GlassFish metrics for DTrace is available. Already known from current versions of WebLogic server is a feature called Active Cache. It's a tight integration of Coherence as an HTTPSession replication mechanism for clustered environments. Another part is the integration with Oracle Access Manager (OAM). The Loadbalancer Plugin is shipped as a ZIP bundle that could be installed separately to GlassFish.

Upgrade Paths

All 2.x installations can use the Upgrade Tool. It replicates the configuration of a previously installed server to the target installation. The process is known as side-by-side install. The Upgrade Tool assists in upgrading the configuration and applications from Sun Java System Application Server 9.1 Update 2 up to and including Oracle GlassFish 3.0.1 to GlassFish Server 3.1. All 3.x installations can make use of the other three update tools (Update Center, Software Update Notifier and pkg cmd-line util). They perform an in-place update of an existing installation. Thinking about the applications, there is hardly any need to change them if they comply to Java EE 6 already. Nevertheless you could be forced to correct some minor problems; especially as the bug fixing efforts by some teams may lead to validation errors or correct lazy interpretations which force you to tweak some of your implementations.

Conclusion

What was promised has been met: Oracle presented GlassFish 3.1 in a timely fashion and with all the announced features. The first, high available Java EE 6 server once again leaves many competitors behind. Commercial and other OSS competitors are on the way but still not there. If you want to use the most current Java EE version on large scale, then you simply don't have an alternative at the moment. And this is not a bad choice in general. The implementation has been tested and stabilised for more than one year now and it has proven its stability many times already. Nevertheless there are some open political issues. The cooperation with non-Oracle suppliers of reference implementations is potentially explosive. Weld, the reference implementation of CDI is a good example. The included 1.1.0 was released on 17th January 2011. The next minor milestone (1.2) is one month behind schedule and will include 95 bugfixes with 3 blocker (one memory related) and eight critical bugs. It's targeted towards JBoss AS 7. Seeing these logical priorities, GlassFish as a reference implementation will always be the first to have 1.0 releases and probably the last to include the more stable x.2 releases.

The initial work on a common platform together with the WebLogic Server has been finished noiselessly and this manifests itself clearly in the significantly better performance in the whole server and cluster. First joint deployment descriptors followed but it's still open to speculation what the final implementation of the common base will look like. The Premier Support for GlassFish 3.1 ends in March 2016. Extended Support is available through March 2019.

> "What was promised has been met: Oracle presented GlassFish 3.1 in a timely fashion and with all the announced features."

This article has been published in Java Tech Journal 3.2011. Read more from Markus in Java Tech Journal 2.2011, where he introduces JBoss Arquillian. More info on www.jaxenter.com

Markus Eisele works at msg systems AG in Munich. He is an Oracle ACE Director for FMW & SOA and UKOUG Member. He runs a blog about Java EE and GlassFish under http://blog.eisele.net and you can follow him on http://twitter.com/myfear.

References

[1] http://weblogs.java.net/blog/sdo/archive/2011/03/01/whats-new-glassfish-v31-performance
[2] http://www.youtube.com/watch?v=O9TCWZ-nIgo
[3] http://blogs.sun.com/gfsecurity/entry/what_s_new_in_glassfish

VOLUNTEER HIGH KICK!

VOLUNTEER DAY 5TH JULY 2011 ...BE THERE!

A date for the calendars of all UKOUG volunteers, Oracle Liaison, staff and Directors of the UKOUG: Tuesday, 5 July 2011 – 2011 UKOUG volunteer day.

The meeting is open to all UKOUG volunteers and Oracle Liaison. The volunteers' day is the UKOUG's annual one day opportunity for all volunteers, Oracle Liaison, staff and Directors to get together.

The primary purpose of the meeting is to learn and discuss the latest in terms of the business plan, strategy and events but, like all great UKOUG events, volunteers' day will involve four key elements:

- Knowledge sharing
- Discussion
- Networking
- Food and drink

Debra, will share her views from her relatively new role as UKOUG Chairman. Margaret Wright, Oracle, Operations Programme Manager and Tom Scheirsen, EMEA User Groups Relationship Manager, Oracle will give us Oracle's views on their involvement in user groups.

Then, to the heart of the day: the 2011 UKOUG Business Plan. A key session as, like all UK businesses and organisations, the UKOUG is affected by the current economic climate.

This will be followed by a lengthy period of group-based discussion of our plans.
Then it's time for lunch and networking, before we discuss "volunteer strategy". Finally, 90 minutes of traditional "community breakout" sessions with the opportunity for each SIG to meet and discuss issues with the Conference committee.

After the business of the day, it's time for part two: Networking, food and a drink. All meeting attendees are invited and encouraged to remain overnight - for those who can spare the time! – for a drinks reception and dinner. A unique opportunity to network with the volunteers, Directors, staff and Oracle Liaison.
info@ukoug.org.uk

UKOUG Annual General Meeting: 5.15-6.00pm. Tucked usefully in between the business of the day, and the evening relaxation, the UKOUG (a not-for-profit company) will hold its AGM, to which all members are cordially invited.
I encourage all members and Oracle Liaison to attend this meeting on Tuesday 05 July as there is, more than in any previous year I can recall, so much to be learnt and discussed as we all endeavour to continue to deliver what is acknowledged as the premier Oracle User Group in the world - in difficult times.

If you would like to stay in touch with Debra, you can follow her on Twitter:

twitter.com/debralilley
twitter.com/UKOUG_chairman

UK OUG
UK ORACLE USER GROUP

2011 DATES FOR YOUR DIARY

UKOUG EVENTS

14-15 June	OUG Conference Series EPM & Hyperion 201 - Mercedes Benz World, Weybridge
15 June	**LOSUG - UKOUG Solaris and Open Solaris SIG Meeting - Oracle City Office**
16 June	UKOUG Application Express (APEX) SIG Meeting - Blythe Valley Park
22 June	**UKOUG Conference Series PeopleSoft 2011 - Park Inn Heathrow, London**
22 June	UKOUG Development SIG Meeting - Oracle City Office
23 June	**OUG Ireland BI SIG Meeting - TBC**
23 June	UKOUG Conference Series PeopleSoft 2011 - Park Inn Heathrow, London
28 June	**UKOUG JD Edwards SIG Meeting - Oracle City Office, London**
30 June	UKOUG Conference Series Siebel & CRM on Demand 2011 - Hilton Reading hotel
5 July	**UKOUG Volunteers Meeting - Hilton Reading hotel**
12 July	UKOUG Primavera SIG Meeting - Aberdeen
14 July	**UKOUG DBMS SIG Meeting - 76 Portland Place**
14 July	UKOUG Local Government Applications SIG Meeting - Oracle City Office
20 July	**LOSUG - UKOUG Solaris and Open Solaris SIG Meeting - Oracle City Office**
17 Aug	LOSUG - UKOUG Solaris and Open Solaris SIG Meeting - Oracle City Office
13 Sept	**UKOUG Apps DBA for Oracle E-Business Suite SIG Meeting - Reading**
13 Sept	UKOUG RAC & HA SIG Meeting - Reading
13 Sept	**UKOUG Security Special Event - Bletchley Park**
14 Sept	OUG Ireland BI SIG Meeting Oracle – Dublin
14 Sept	**OUG Ireland HCM SIG Meeting Oracle - Dublin**
14 Sept	UKOUG Application Server & Middleware SIG Meeting - Thames Valley Park
15 Sept	**UKOUG Oracle Financials SIG Meeting - London, TBA**
20 Sept	UKOUG Development SIG Meeting - Oracle, Thames Valley Park
20 Sept	**UKOUG Management & Infrastructure SIG Meeting - Oracle City Office**
21 Sept	LOSUG - UKOUG Solaris and Open Solaris SIG Meeting Oracle City Office
22 Sept	**UKOUG Public Sector HCM Customer Forum - Blythe Valley Park**
22 Sept	UKOUG UNIX SIG Meeting - Thames Valley Park
28 Sept	**.NET on Oracle on Windows SIG Meeting - Oracle City Office**
29 Sept	UKOUG Oracle Projects SIG Meeting - Oracle City Office
4 Oct	**UKOUG Education & Research SIG Meeting - Blythe Valley Park**
11 Oct	UKOUG HCM SIG Meeting - Maple House
11 Oct	**UKOUG Scottish Development SIG Meeting - Sun Oracle, Linlithgow**
12 Oct	UKOUG Business Intelligence & Reporting Tools SIG Meeting - RIBA
12 Oct	**UKOUG Hyperion Planning & Essbase SIG Meeting - Oracle City Office**
13 Oct	UKOUG Public Sector Combined Event - CBI Conference Centre
18 Oct	**UKOUG DBMS SIG Meeting - Blythe Valley Park**
18 Oct	UKOUG Hyperion HFM SIG Meeting - Oracle City Office
19 Oct	**LOSUG - UKOUG Solaris and Open Solaris SIG Meeting - Oracle City Office**
20 Oct	UKOUG Business Intelligence & Reporting Tools SIG Meeting - RIBA
20 Oct	**UKOUG Local Government CRM Customer Forum - Blythe Valley Park**
3 Nov	UKOUG Application Express (APEX) SIG Meeting - Oracle City Office
8 Nov	**OUG Ireland BI SIG Meeting Oracle – Dublin**
9 Nov	UKOUG Conference Series JD Edwards 2011 - Location TBC
16 Nov	**LOSUG - UKOUG Solaris and Open Solaris SIG Meeting - Oracle City Office**
5-7 Dec	UK Oracle User Group Conference 2011 - Birmingham ICC
21 Dec	**LOSUG - UKOUG Solaris and Open Solaris SIG Meeting - Oracle City Office**

UK Oracle User Group, User Group House, 591-593 Kingston Road, Wimbledon, London, SW20 8SA

Tel: +44 (0) 20 8545 9670 Fax: +44 (0) 870 9000 335 Email: info@ukoug.org Web: www.ukorg.org

E-BUSINESS SUITE

PZ CUSSONS IMPROVES ITS DATA ACCURACY AND KNOWLEDGE SHARING WITH ORGPLUS ENTERPRISE

Gill Daniels describes how an HR team responsible for a global organisation can use HR charting software, integrated with Oracle HR, to display up the minute data to staff worldwide, improving HR data accuracy and knowledge sharing. On a higher level, managers can also view HR data, sharing best practice on departmental structures across the organisation. With additional interactive modelling capability, an HR charting tool can also be used to support strategic HR planning, aiding mergers and acquisitions and succession planning.

Large organisations with global reach that have specific local brands, often face the challenge of sharing data across country boundaries to maximise the management of human capital. Managing a widespread organisation requires investment in central IT systems – HR, financial or operational that can be effectively shared across the regions or countries.

An example of this is PZ Cussons, which operates in Africa, Asia and Europe. The HR team at PZ Cussons has invested in an Oracle On Demand HR system to manage its HR data for all personnel worldwide. The HR system provides an easy to use central repository accessible by all of the HR services teams. OrgPlus Enterprise charting software has been integrated with Oracle On Demand, helping PZ Cussons to maximise on its investment in the Oracle system.

OrgPlus enables an organisation to be displayed in visual charts, not just for the managers and the HR team, but also to provide an overview for employees across the organisation, providing transparency and visibility. This access to their own HR information for all employees has resulted in a dramatic reduction in the number 'orphan' HR records in the Oracle system from 20% down to just 2%.

PZ Cussons initially started using OrgPlus Enterprise in the UK and has since rolled it out across its locations in Nigeria and Indonesia. All employees are able to view the organisational charts via the intranet – those that are not full time PC users are provided with secure login details so that they can use kiosks, available for all staff to use.

By drawing data from Oracle, OrgPlus Enterprise enables PZ Cussons HR managers to validate data accuracy and ensure that the correct staff member is shown as reporting to the right person in the organisation. The charts help the HR team to identify 'orphans' where the data shows that individuals are not recorded as reporting to anyone.

As well as position in the organisation, contact email addresses are displayed and provide an immediate link for employees to connect up with others.

Planning Tool for HR and Managers

HR teams can also use OrgPlus Enterprise to validate and capture accurate HR data and make it widely available to managers, wherever they are located. Managers can access head count information and department structures that help with resourcing and planning. They can view other departments, using knowledge sharing for adopting best practice on team structures and positions.

Supporting Strategic Planning

Organisations like PZ Cussons that are already reaping the benefits of charting software at a practical level are now taking it one step further. Blueprint, an interactive modelling module for OrgPlus Enterprise helps with planning acquisitions and merger activity and is being implemented by PZ Cussons. With Blueprint the managers can perform organisational design exercises and see how change options would impact budgets and reporting structures.

Because it is web based, it enables the managers across different locations to share ideas and make informed decisions to shape the new organisation around current needs, as well as plan for future growth.

We have also found that some HR teams, like that at PZ Cussons, plan to use OrgPlus to support decision making at a strategic level, to identify key roles and develop succession and talent management plans to address business continuity risks.

UKOUG AUTHOR PROFILE

Gill Daniels is Senior Director, EMEA, at HumanConcepts has more than 20 years' experience with enterprise technology solution and service providers to blue chip organisations worldwide.
Contact:
gillian.daniels@humanconcepts.com
Human Concepts
Tel: 01392 829150

HYPERION & ORACLE

by Julie Harris

It's been more than three years since Hyperion was acquired by Oracle to extend and strengthen its application and technology offerings for Corporate Performance Management and Business Intelligence.

Much has been achieved in this time including the addition of new modules including:

- Financial Close Manager and Disclosure Management;
- Functional enhancement of existing modules such as Hyperion Planning and Financial Data Quality Manager; and,
- Increased product integration to the broader Oracle application and technology portfolio.

If anyone ever questioned the future of Hyperion and its applications and products, Oracle has firmly demonstrated that this portfolio is here to stay!

At a time when we have seen unprecedented and continuous change in the UK and global economy it reminds me of the famous statement from Winston Churchill "He who fails to plan is planning to fail" – a statement that still holds true today, perhaps more than ever before.

There is no doubt that organisations small and large are embracing this principal in its fullest. They are extending the value they get from technology systems that give them their own unique insight into their business allowing them to proactively respond to change rather than be a victim of change.

The article from Mehmet Hamit in this edition of Oracle Scene considers this principle in the context of how Hyperion Planning addresses the challenge of building and maintaining a plan that allows you to adapt and manage change effectively.

Whether it's better planning, better reporting and/or better intelligence, the Oracle Hyperion applications and solutions are there to provide additional value and unlike some technologies can be adopted and implemented quickly and effectively without the need to rip and replace large parts of any existing IT landscape (including a non-Oracle one!). So what should an Oracle User be doing now? Whether you are IT, an end user or in a line of business, I bet you receive, manage or even create spreadsheets containing corporate information and intelligence! I also bet you have moments of frustration when you just can't get the information you need to do your job and make important business decisions!

On the 14th - 15th June 2011, OUG Conference Series EPM & Hyperion 2011 is providing the perfect opportunity for more than three hundred users, who share similar problems, to discuss hints and tips and help challenge thinking, at Mercedes-Benz World in Weybridge.

UKOUG are delighted to announce that this year's event will, for the 1st time, include FREE Hands-On Workshops; allowing users to touch, feel and assimilate some of the products in guided sessions.

In attendance will be: Oracle Hyperion and Oracle Business Intelligence, Oracle ACE's, Oracle Hyperion Customers and Partners.

If you would like to hear further information on EPM & Hyperion please contact: hyperion@ukoug.org.

Alternatively we are excited to announce related SIG events on the following dates:

UKOUG Hyperion Planning & Essbase SIG Meeting
Date: Wednesday 12th Oct 2011
Location: London

UKOUG Hyperion HFM SIG Meeting
Date: Tuesday 18th Oct 2011
Location: London

UKOUG AUTHOR PROFILE

Julie Harris is currently Oracle UK EPM Business Development Manager and the Oracle Liaison for the Hyperion User Community. Julie has over 15 years experience in Performance Management, Business Intelligence and Data Warehousing with leading vendors Oracle/Hyperion, IBM and Teradata, preceded by Project Manager and IT Technical roles with Lloyds TSB and Municipal Mutual Insurance (since acquired by Zurich Insurance).

HYPERION

PLANNING FOR HYPERION PLANNING

By **Mehmet Hamit**
Manager, International Financial Reporting System, Henry Schein

Hyperion Planning is a part of the Oracle Enterprise Performance Management suite, and is built on top of Oracle Essbase, part of the Oracle Business Intelligence suite.
It is a centralised Excel and Web-based planning, budgeting and forecasting solution.
Hyperion Planning or Hyperion has been a budgeting solution in some form for a long time; firstly through Hyperion Pillar in 1994, the biggest customer was BT and Dell. Now Hyperion is a part of Oracle, and rebranded as Planning, it is still the clear market leader in the planning and budgeting market with its main competitors.
Planning facilitates the relationship between people and process throughout the budgeting cycle, by encouraging a "closed-loop" approach:

- Setting objectives and goals;
- Allowing advance modelling;
- Ability to track and audit relationships; and,
- Analysis and reporting.

The technology interfaces with Microsoft Office, is Web based and is centrally administered, thus Planning provides the core technology required to establish inflection of control and predictability to the budgeting and forecasting process.

The Business Case for Planning

Typically in applying a Planning solution the capabilities would allow for several improvements including:

- Unique and identifiable spreading based on previous years;
- A single form with current year, prior year-actual and current year budget;
- Cost of sales based on segments and products of the business;
- Automated overhead allocation;
- Planning at employee level on drivers such as taxes and benefit;
- Planning based in fiscal and tax calendars;
- Employee benefit in kind planning;
- Payroll allocations; and,
- Incorporation of employee Travel and Entertainment expenses.

Capital Planning

In addition to the above, we also require the following:

- Custom forms for use with budgeting, standard assets, leasehold improvements and capital projects;
- Accumulation of existing and new asset depreciation information;
- Asset lifecycles based on predefined assets classes;
- Supporting detail for assets within each asset class; and,
- Integrated acquisition and depreciation information for planning.

The benefits of implementation to support a business case are numerous, but include:

- A more consistent budgeting process;
- One system for all budgeting needs;
- Easier reporting systems;
- Accurate, timely and up-to-date allocations;
- Elimination of dual entries;
- Faster and more consistent consolidation;
- Allowance for more budget users;
- Improved source documentation within the budgeting toll;
- Better control of the budgeting process with use of workflow; and,
- Alternative organization budgeting structures.

Hyperion Planning/Spreadsheets

Spreadsheets are handy for calculating a row of numbers, and they are on the whole helpful if you want to do calculations on-the-fly, and even automate the planning process using VBA. As an accountant, where we start to have difficulty though, is when you want to distribute the planning process amongst managers and Directors within the business, as the planning process goes through much iteration, it becomes difficult to orchestrate and control your working due to the number of spreadsheets being populated. By having Planning as your central shared planning application you can model scenarios, allocate costs within various hierarchies, and then roll them up within the rest of the financial planning process, simply.

The Business Partner

There is no shortage of business partners or business intelligence vendors out there. They all have a track record with easy-to-use, flexible and affordable solutions. So how to you pick the right partner? Here are a few tips to be ahead of the curve when typically choosing a BI solution partner. Visualization of data is important of course, but the biggest mistake you can make is judge the BI vendor based on the pretty dashboard samples they show you on their website or during a demo. Every BI vendor can do that because visualization software components are cheap. The real challenge is customizing these dashboards to your own needs and having them show your own data. This part usually takes most vendors months, and costs you a small fortune. If the BI vendor cannot get your own data to show the way you like it within just a few days, move on.

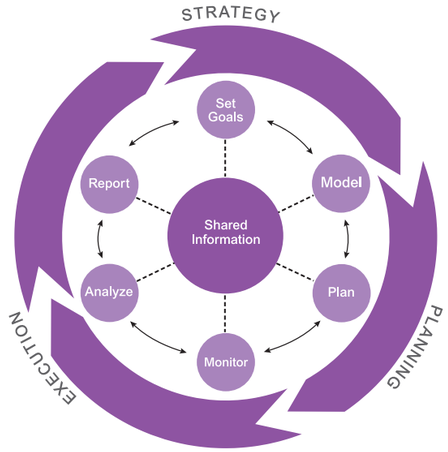

Data Warehouse

Today's BI technology does not require a data warehouse, even when there are multiple data sources involved, large amounts of data or multiple users querying the data. There are very specific scenarios where a data warehouse is a good idea, but they are most likely not relevant to you. If the vendor requires a data warehouse to proceed with implementation, progress on.

OLAP

OLAP is also very lengthy and costly to implement, and there is really no need for it anymore. Today's BI technology can handle even huge amounts of data without OLAP, at fractions of the time or cost. If the BI vendor requires OLAP to assure you acceptable query performance. Move on.

Significant Upfront Investment

Do not agree to this, and demand to have at least one solid report or dashboard running over your own data before you commit to anything significant in advance. If the vendor is not willing to do so, it's probably because they would have to spend weeks on development before they can reach that point. That typically means this vendor is either using very old technology or is simply trying to pull one over you.

Professional Services

If you choose a BI vendor who makes most of his business from professional services, you can pretty much be sure that they will take their time building your solution. These types of BI vendors also live off on-going maintenance services, so what you initially pay for the solution is actually only the beginning. Whenever possible, try to choose a BI vendor that focuses on selling BI software to the end customer, not to the professional services community.
Finally, the most important thing is to get the vendor to prove what they claim prior to investing too much money upfront. This proof must be in the form of reports, dashboards or analytics in real life scenarios, running on real data, used by the actual end users and within a reasonable amount of time. If a vendor is not willing to accommodate this simple request, you really should find one that does. Many vendors provide free trial versions, as well as utilise technology that speeds up implementation tremendously. If the one you're in contact with now doesn't, move on.

PEOPLESOFT

The Past Present & Future of PeopleSoft

By Marc Weintraub, Oracle's Director of Product Management and Development

The Past Delivery of PeopleSoft 9.1

When looking at where PeopleSoft is today and where it is going tomorrow, one should start in the past. The past here is the recent past, the fall of 2009 to be precise. That is when Oracle unveiled the latest release of Oracle's PeopleSoft, Release 9.1. PeopleSoft 9.1 is one of the most robust and comprehensive releases in PeopleSoft's history. It includes 21 new solutions, 1,350 new features, more than 28,000 pages enhanced with Web 2.0 capabilities, 300 new Web services, and 200 industry-specific enhancements. This release contains significant enhancements across all of the major PeopleSoft product areas including Human Capital Management, Financials, Supplier Relationship Management (Procurement), Enterprise Service Automation (Projects), Supply Chain and Asset Life Cycle Management solutions.

The Present

Here in the present, a bit more than one and a half years after the release of PeopleSoft 9.1, adoption of this release is three times faster than previous PeopleSoft releases across all products. Over 1,250 customers are live on or deploying PeopleSoft 9.1. Some of the customers recognizing the value of PeopleSoft 9.1 and who have implemented or are implementing the latest release, including The Dow Chemical Company, Papa John's, the Defence Science and Technology Laboratory (Dstl) part of the UK Ministry of Defence (MOD), Portland General Electric, China Resources (Holdings) Co. Ltd., Premera Blue Cross and Husky Injection Molding Systems Ltd. In addition, the PeopleSoft community is expanding with more than 350 new customers selecting PeopleSoft applications over the last 12 months. They include: the Carlyle Group, BJC HealthCare, Navigant Consulting and Amedisys.

Feature Packs

The shifting demands of PeopleSoft customers for more capabilities has provided an opportunity for Oracle to adjust its investment strategy for PeopleSoft. Oracle's planned investment strategy for PeopleSoft is to value-add capabilities while at the same time building and delivering major new releases. The release methods will be Feature Packs:

- New PeopleTools Releases
- New Major Application Releases

To do this, additional enhancements are planned to be delivered between major releases onto the current PeopleSoft release as part of regular maintenance. Oracle plans to roll up all previously delivered off-cycle enhancements and maintenance annually into a Feature Pack. Feature Packs include a recut CD and include updated documentation and recertified upgrade scripts. Oracle also plans annual PeopleTools releases as well as major releases for PeopleSoft applications every three years.

The already delivered Feature Packs for PeopleSoft (PeopleSoft HCM 9.1 Feature Pack 2010 PeopleSoft FMS/SCM 9.1 Feature Pack 2011) include previously delivered off-cycle enhancements and offer significant new features and capabilities to PeopleSoft 9.1 customers prior to their inclusion in PeopleSoft 9.2.

The PeopleSoft HCM 9.1 Feature Pack 2010 includes the following capabilities:

- Company Directory provides configurable organisation visualisation and navigation capabilities that allow end users to create a personalised view of their company's user directory;
- Oracle's PeopleSoft Talent Management 9.1 to PeopleSoft HCM Integrations that connect PeopleSoft Talent Management 9.1 applications and earlier versions of PeopleSoft Human Resources applications;
- Clairvia Healthcare Scheduling to HCM integration is available through a partnership with Clairvia;
- Flexibly configure different types of leave rules for Country, Federal, State, Local, and company-defined programs within the Extended Absence/Leave functionality; and,
- Crystal Reports were rewritten using Oracle's XML Publisher (BI Publisher).

New capabilities in PeopleSoft FMS/SCM 9.1 Feature Pack 2011 include:

- PeopleSoft Mobile Inventory Management, a new application within Oracle's PeopleSoft Enterprise Supply Chain Management suite;
- Enhancements to Oracle's PeopleSoft Strategic Sourcing module with capabilities around managing documents associated with RFQs and Bids;
- Important integrations that streamline the period end close process. Oracle's Hyperion Disclosure Management is now integrated with PeopleSoft Financials, allowing customers to support XBRL financial statement reporting to regulatory bodies. Additionally, Oracle's Hyperion Planning is now integrated to PeopleSoft Financials 9.1 enabling customers to plan and manage organisational budgets at all levels;
- Extended automation of financial controls;
- Enhanced support for global deployments for Europe, Asia and North America;
- New web services following HR-XML standards;
- Accelerated financial processing of massive project data volumes and complex revenue contracts; and,
- New dynamic bonus depreciation that enables organisations to lower the cost and effort of adapting to new legislation as soon as it is enacted.

The Future

Whether it is included in a Feature Pack or part of the next PeopleSoft release, Oracle's planned product road map for PeopleSoft applications is to deliver value-add features for all of an organization's constituents along three design principles—Simplicity, Productivity, and Lowered Total Cost of Ownership (TCO).

Delivering greater simplicity will enable casual users to effectively leverage PeopleSoft applications while remaining focused on their primary job function. By delivering rich features and a flexible role-based user experience, we continue to enable business owners to improve productivity and execute mission-critical business transactions. Finally, for the IT department the TCO is lowered by adding new functionality that allows organisations to eliminate costly customisations and by expanded integrations within PeopleSoft and between PeopleSoft and other Oracle applications (for example, Hyperion, GRC, and Primavera). Productivity, and Lowered TCO are expected to remain at the core of Oracle's planned capabilities for PeopleSoft applications for the foreseeable future. Oracle intends to regularly validate these design principles and its investment strategy for PeopleSoft with its customers through Customer Advisory Boards, Strategy Councils, and direct one-to-one customer engagements to ensure that Oracle delivers the most long-term value as possible to its customers.

UKOUG AUTHOR PROFILE

Marc Weintraub is a Director of Product Management and Development for PeopleSoft Enterprise applications within Oracle's product development organisation. Marc provides content and works directly for Paco Aubrejuan, the Vice President and General Manager of the PeopleSoft Enterprise division and works closely with John Webb, the product management Vice President and his entire team. Since starting his role in mid 2006 Marc has evangelised the PeopleSoft Enterprise solution by delivering product line overviews and updates to over 7,500 customers at more than 500 customer facing events.

Marc Weintraub has 16 years of experience in the software industry, including the past 6 years at PeopleSoft/Oracle. He has experience multiple roles in the past including: product marketing, product management, technical pre-sales, and systems implementations. Marc joined Oracle from PeopleSoft in 2005 where he focused on marketing Oracle's procurement solutions.

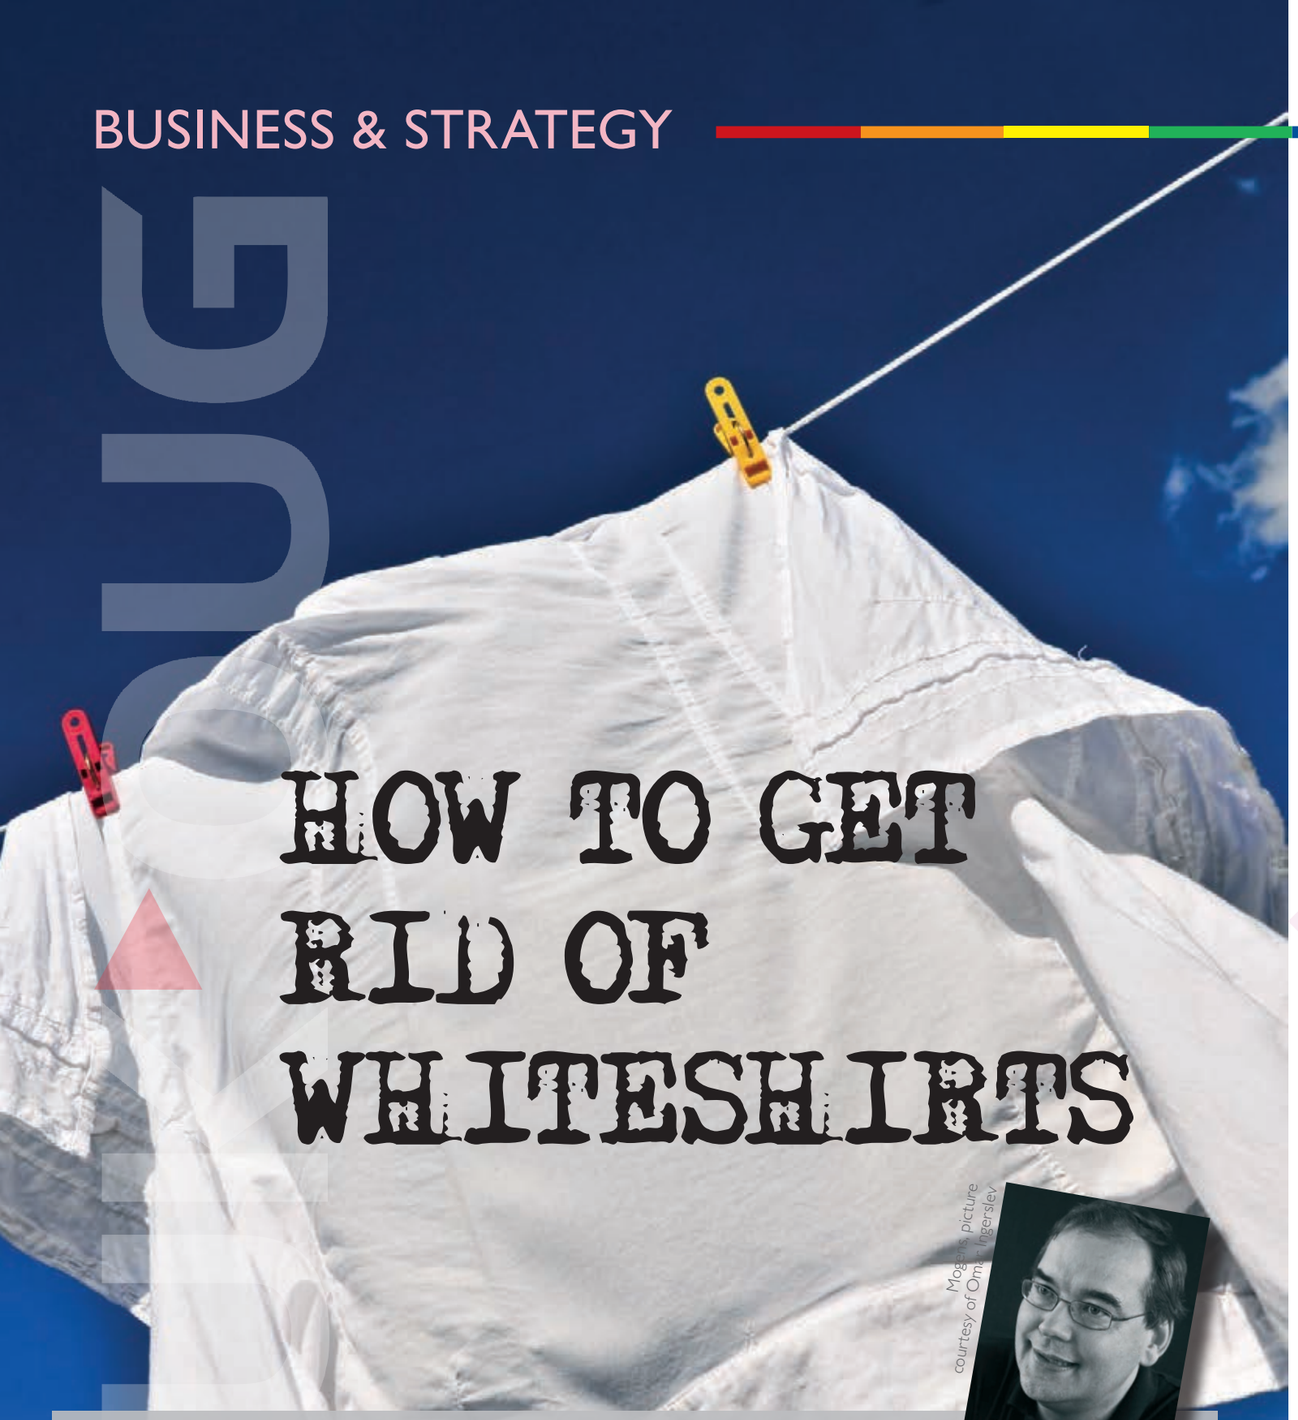

HOW TO GET RID OF WHITESHIRTS

Mogens picture courtesy of Oma Ingerslev

As I have tried to explain to you, my dear friends in the UK, life has become much less pleasurable for me and my trusted database guru nerd expert superheroes here at Miracle A/S.

First, my co-CEO and former friend Lasse started wearing shark-skin suits and blue suede shoes and the like. That should have warned me. I just thought he'd found a lustrous mistress. But it was worse: He'd found the Oracle E-Business Suite, which to a certain extent can bring down much bigger empires than even the most beautiful mistress (and much harder, if I may say so).

Then he slowly started to change our flat, no-frills, well-working, super-profitable organisation into the opposite, namely one that was suited (!) for delivering OEBS implementations. Changes to the organisation here at Miracle A/S include, but are in no way limited to, the following:

- An extra layer, ie. middle managers, in order to slow down and obfuscate communications, resulting in longer projects (hence more Kroner to us);
- Lots of new procedures w.r.t. billing of customers and payment of bills from our vendors. These include mandatory procedures for losing bills, getting confused over simple bills that are very clear, and extending the payment period to 420 days. Not because we like it, but in order to get a feel for those of our customers who run SAP or OEBS;
- Dress code 25/8/370 should now include skirts (so far only for female staff), white shirts (for everyone), oily hair, and stockings (men can wear black socks instead, as long as they have the same length as stockings, so no white legs are ever seen);
- A couple of long, internal memos from Lasse about how to park correctly in the parking lot, and how to minimise time in the bathrooms, both in order to increase number of billable hours for the consultants;
- No more free lunches, but a mandatory health program named "Fat for Life", which includes monthly reports to Lasse from all of us detailing how many metres we have walked, run, and crawled per day (in a spreadsheet, of course);
- A new company car policy called "Company Car 2020 Plus", which effectively means no more company cars to anyone except Lasse.

I should of course have seen it coming. But I don't think anyone can predict the impact of the Tsunami of software: the ERP-programs. Man, can they move some stuff around and place an organisation on its head! So life is getting more and more miserable. I'm receiving urgent calls from bereaved wives of former men, telling me of sleepless nights and endless ironing, of husbands begging on the streets for a cup of gasoline, and worse.

But in the middle of this, just when we were looking into the very heart of darkness, along came our annual conference, the Miracle Open World 2011 conference at Legoland.

Lasse, me, the slaves, the customers, and some partners were all there, and since the majority of the attendees are non-Miracle (and since the music in the party house was loud), there was no way Lasse could force his OEBS regime on them. So:

- We had open and happy discussions about new SQL features in 11gR2;
- We arranged empty beer bottles on the table (as we emptied them) so as to represent every guardsman in the Royal Guards formation walking through Copenhagen to the royal castle (one of the nights we had to include the accompanying big-band, too);
- We had our beach party at the local water park, including four-person rafts;
- We had sessions, breaks, lunches, dinners (and after-sessions in the party house) where we could freely talk to Graham Wood, Jonathan Lewis, Cary Millsap, and all the other true heroes that haven't succumbed to OEBS yet;
- We could even – but don't tell Lasse this – crack a joke or two about ERP-systems, whiteshirts, process managers, programME managers (as opposed to, one must emphasise sharply and sternly, program managers) and bills not getting paid on time anywhere anymore thanks to the streamlining of the payment processes (whether it be upstream, downsteam, or downdrain).

Somehow, it helped a lot. It was nice to speak freely again like free men, and it dawned on me that I had finally – after all these years – found one good reason to go to conferences.

It also explained why all these whiteshirts are present at booths during, say, the big UKOUG conference, in Birmingham, although nobody wants to talk to them and they don't sell anything:

They get away from the local OEBS or SAP regime, if only for a few days.

PS: Lasse is talking about implementing OEBS for our brewery, since we've had problems with the machine that puts beer into kegs. He says we need to streamline the processes in order to prevent future problems, and in order not to depend on certain specialised staff.

Debra Lilley adds - "Rubbish Mogens!!, I even talked about Apps at YOUR conference. There is no need for a database without an application"

Mogens Nørgaard

SUPPORT ZONE

ORACLE CONFIGURATION MANAGER (OCM)

An overview by Kate Cumner, Oracle Support Services

Oracle Support provides proactive system health checks, patch advice, and other valuable information about your Oracle products. It is part of your Oracle Premier Support service.

Oracle Configuration Manager (OCM) is the key to enabling these proactive capabilities of My Oracle Support, along with priority handling of Service Requests and configuration-driven personalisation.

The benefits seen by OCM users include improved uptime:

- 25% less internal support time on Database Administration;
- 50% faster issue resolution;
- 25% problem avoidance with Alerts and Health Check;
- 30% reduction in the time it takes to log a Service Request; and,
- 80% reduction in the time taken to apply critical patch updates.

OCM automatically collects information on your configuration information. This enables My Oracle Support to provide specific health & recommendations, patch plans, inventory reports, coupled with the configuration and change management capabilities.

Collected data is sent securely to Oracle Support. The result is:

- Faster problem resolution by integrating your configuration information into the service request flow providing Oracle Support the information they need real-time to resolve your problem quickly and efficiently;
- Improved systems stability delivered through proactive advice & health checks driven by Oracle best practices and personalised to your system configuration; and,
- Simplified configuration management from a single, comprehensive and personalised dashboard of configurations, and inventory reports.

Where do I start?

1. Login to My Oracle Support (http://support.oracle.com)
2. Open the Collector Tab
3. Review the Quick Start Guide

This gives the key steps:

- How to identify if OCM is already installed.
- Downloading OCM.
- Installing of OCM.

OCM can be configured in Connected Mode or Disconnected Mode. The Guide also advises on Instrumenting Oracle Database instances, and associating configurations with a Customer Support Identifier.

Introducing OCheck

OCheck was first introduced in OCM 10.3.3. It provides the ability to evaluate diagnostic information associated with an Oracle environment, enabling a rich set of precise and targeted My Oracle Support health recommendations. OCheck does not collect or upload any diagnostic data to Oracle; it merely checks for a condition and uploads a true / false result as part of the standard OCM dataupload process. The My Oracle Support health check rule logic back at Oracle leverages these results, along with the uploaded configuration data, to determine the health of the customer's system. The subset of health checks that require OCheck to be configured are clearly identified in the "Pre-Req" column of the My Oracle Support Health Checks Catalogue [ID 868955.1]

Frequently Asked Questions

What Information does Oracle collect? What about the Security & Privacy of my information?

OCM collects only the configuration information that is typically requested by Oracle Support when opening a service request. The OCM Collection Overview document (MOS note 728985.5) lists all the information that is collected by OCM.

Collected information includes:

- Patch levels.
- Deployed components and respective versions, and type
- Content of Configuration files.
- No sensitive data such as business transactions, production data, or passwords is collected by OCM.
- OCM has to be enabled in your environment before it collects data.
- OCM configuration information is collected as XML files and transmitted to Oracle using SSL encryption.
- OCM can work in "disconnected mode" allowing you to review the content of the XML files before you send them to Oracle.
- The Oracle Support Hub allows many OCM collectors to use a single internal point to upload their configuration data, eliminating the need for each individual OCM collector to access the Internet.
- All collected information is secured in Oracle's environment in accordance with Oracle's security support policies.
- For full security policy information please refer to Oracle Global Support Security (http://www.oracle.com/us/support/library/customer-support-security-practices-069170.pdf).

Doesn't Oracle Enterprise Manager provide these functions? Do I need both OCM and OEM? – Will this create overheads?

- Oracle Enterprise Manager (OEM) stores information centrally in your environment, enabling you to monitor, configure and maintain Oracle products.
- OCM sends information to My Oracle Support, and is used by Oracle Support for diagnosing issues and provide proactive alerts and health checks.
- Customers using Enterprise Manager Grid Control 10.2.05 or greater can use EM Harvester to centralise the collection and uploading of configuration data, for My Oracle Support.
- OEM 11g adds further integration with My Oracle Support by providing Knowledge Browse & Search, Service Requests, Proactive Patch Recommendations and Patch Plans capabilities.
- Customers do not need to install OCM if they are using Enterprise Manager, they just need to enable EM Harvester on the OMS Repository.
- OCM defaults to send information daily but can be configured to update weekly or other intervals; and using the EM Harvester allows OCM to send a single file, reducing any overheads still further.

Do I need to be connected? What options are there for collecting configuration data?

- Oracle Configuration Manager installs in each Oracle home. Each OCM instance can then connect directly with Oracle's support systems or connect through a proxy removing the need to open systems' ports for connectivity to the Internet.
- For larger deployments, Oracle provides in the Companion Distribution a set of utilities to centrally manage the deployment and configuration of OCM.
- Enterprise Manager customers do not need to install OCM. Configuration data is collected using the Harvester.
- OCM can be used in disconnected mode to collect and upload configurations when a direct internet connection from your host machine to Oracle is not possible. The local files can be reviewed before being manually attached to Service Requests – [see note 453412.1].

Where can I find more information?

Installation and Admin Guide - http://download.oracle.com/docs/cd/E17344_01/doc.103/e16665.pdf

Testing Network Connectivity – http://download.oracle.com/docs/cd/E12482_01/doc.103/e12882.pdf

OCM Detailed FAQ - https://support.oracle.com/CSP/main/article?cmd=show&id=548847.1&type=NOT

Configuration Manager Documentation - http://download.oracle.com/docs/cd/E18041_01/index.htm

Pre-requisites – MOS note id 728473.5

My Oracle Support notes 603505.1 and 418277.1 have further details on recorded training and scheduled webcasts -

- Installation and Administration Guide
- Pre-requisites
- Testing Network Connectivity
- Getting Started with OCM
- OCM Detailed FAQ
- Knowledge Articles about OCM

UKOUG

UK ORACLE USER GROUP

Benefits of UKOUG membership:

- Free places at UKOUG events*
- 10% discount on Oracle University courses in the UK and other training providers
- Up to 40% off Oracle titles from leading publishers
- Access to digital and print editions of UKOUG publication, Oracle Scene
- Regular ebulletins delivering the latest community specific news, events and information
- A range of online resources, including an extensive library of conference papers, event presentations, articles, survey results and more
- The opportunity to nominate yourself for key roles within an established community
- Hearing first hand about Oracle developments and being able to collectively influence future direction

**Dependent on level of membership*

For more information on membership benefits:
UK Oracle User Group | 591-593 Kingston Road | Wimbledon | London | SW20 8SA
T: +44 (0)20 8454 9670 F: +44 (0)870 9000 335 Web: www.ukoug.org

Serving the Oracle Community


UKOUG
PARTNER OF THE YEAR
AWARD WINNER 2011/2012
Together...
We made it
Possible!
18+ TOP
INDUSTRY
AWARDS
E-Business Suite Partner
Hyperion Partner
Siebel Partner
PeopleSoft Partner
JD Edwards Partner
ISV Partner
Database Partner
Middleware Partner
Business Intelligence Partner
Innovation Partner
Training Partner
Managed Services Partner
Public Sector Partner
Financial Services Partner
Retail, Manufacturing and Distribution Partner
Communications, Media & Utilities Partner
Services, Transport & Logistics Partner
Small Business/General Business Partner
Nominations now open:
www.ukoug.org/our-partners/ukoug-pya
UKOUG
UK ORACLE USER GROUP